

Module 1 : Foundations of Bayesian Thinking for Trading

Course: Bayesian Regression and Time Series Forecasting for Commodities Trading

Learning Objectives

By the end of this module, you will be able to:

- 1 . **Understand** the difference between Bayesian and Frequentist approaches in a financial context
 - 2 . **Apply** Bayes' theorem to calculate updated beliefs about market conditions
 - 3 . **Interpret** probability as a degree of belief about uncertain market outcomes
 - 4 . **Implement** basic Bayesian calculations using Python
 - 5 . **Recognize** when Bayesian methods are preferable for trading applications
-

Why This Matters for Trading

Markets are fundamentally uncertain. Traditional statistical methods often assume we have unlimited data and that parameters are fixed constants waiting to be discovered. But in commodities trading:

- **Regimes change:** The dynamics of crude oil in 2020 are different from 2010.
- **Data is limited:** We may only have 50 years of wheat prices, but wheat has been traded for millennia.
- **Uncertainty matters:** A forecast of "\$80/barrel" is useless without knowing how confident we are.
- **Prior knowledge exists:** Expert traders have beliefs about seasonality, mean reversion, and demand.

Bayesian methods let us:

- **Quantify uncertainty** in every prediction
 - **Incorporate domain expertise** through prior distributions
 - **Update beliefs systematically** as new data arrives
 - **Make decisions** that account for model uncertainty
-

1 . The Two Philosophies of Probability

1 . 1 Frequentist View

In the **frequentist** view, probability represents the long-run frequency of events:

- "There's a 60% chance of rain" means "In similar atmospheric conditions, it rains 60% of the time"
- Parameters (like true average return) are **fixed but unknown constants**

- Data are random samples from a population
- We use data to estimate the fixed parameters

The problem for trading: We can't repeat the 2008 financial crisis. Each market is unique. What does "probability of a recession" mean in frequentist terms?

1.2 Bayesian View

In the **Bayesian** view, probability represents our **degree of belief**:

- "There's a 60% chance of rain" means "Given my current information, I believe there's 60% chance of rain"
- Parameters are **random variables** with probability distributions
- We start with **prior beliefs** and update them with data
- The result is a **posterior distribution** reflecting updated beliefs

The advantage for trading: We can talk about "the probability that crude oil will exceed \$100 though this specific future hasn't happened yet."

2. Bayes' Theorem: The Foundation

At the heart of Bayesian inference is **Bayes' theorem**:

$$P(\theta | \text{Data}) = \frac{P(\text{Data} | \theta) \cdot P(\theta)}{P(\text{Data})}$$

Or in words:

$$\text{Posterior} = \frac{\text{Likelihood} \times \text{Prior}}{\text{Evidence}}$$

Let's break this down:

| Term | Symbol | Trading Interpretation | Prior | Posterior |
|---------------------------------------------------------------------------------------------|-------------|-------------------------------------------------------|-------|-----------|
| about parameter θ before seeing data (e.g., "Gold tends to rise during uncertainty") | $P(\theta)$ | | | |
| $P(\text{Data} \theta)$ | | How likely is the observed data given our hypothesis? | | |
| $P(\text{Data})$ | | Total probability of the data (normalizing constant) | | |
| Updated belief after seeing the data | | | | |

The key insight: **The posterior combines prior knowledge with observed data.**

```
In [ ]: # Setup: Import libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy import stats
import warnings
warnings.filterwarnings('ignore')

# Set random seed for reproducibility
np.random.seed(42)

# Plotting style
plt.style.use('seaborn-v0_8-whitegrid')
plt.rcParams['figure.figsize'] = (12, 6)
```

```
plt.rcParams['font.size'] = 12  
  
print("Libraries loaded successfully!")
```

3 . Trading Example: Is This Strategy Profitable?

Let's apply Bayes' theorem to a practical trading question.

Scenario

You've developed a trading strategy and backtested it for 20 days. It was profitable on out of 20 days (70% win rate).

Question: What's the probability that this strategy has a true win rate above 50%?

Frequentist Approach

A frequentist would calculate a confidence interval:

- Point estimate: 70%
- 95% CI: approximately [46%, 88%] (using normal approximation)

But the frequentist **cannot say** "there's an X% probability the true win rate is above 50%". We can only say "if we repeated this experiment many times, 95% of intervals would contain the true parameter."

Bayesian Approach

We can directly answer: **What's the probability the true win rate > 50%?**

```
In [ ]: def bayesian_win_rate_analysis(wins, total, prior_alpha=1, prior_beta=1):  
    """  
        Analyze win rate using Bayesian inference with Beta-Binomial model.  
  
        Parameters:  
        -----  
        wins : int  
            Number of winning trades  
        total : int  
            Total number of trades  
        prior_alpha, prior_beta : float  
            Beta distribution parameters for prior  
            alpha=1, beta=1 gives uniform prior (no prior belief)  
  
        Returns:  
        -----  
        dict with posterior statistics  
    """  
    # Posterior parameters (Beta-Binomial conjugate update)  
    posterior_alpha = prior_alpha + wins  
    posterior_beta = prior_beta + (total - wins)  
  
    # Create distributions  
    prior = stats.beta(prior_alpha, prior_beta)
```

```

posterior = stats.beta(posterior_alpha, posterior_beta)

# Calculate key quantities
results = {
    'posterior_mean': posterior.mean(),
    'posterior_std': posterior.std(),
    'prob_above_50': 1 - posterior.cdf(0.5), # P(win_rate > 0.5)
    'prob_above_60': 1 - posterior.cdf(0.6), # P(win_rate > 0.6)
    'hdi_95': (posterior.ppf(0.025), posterior.ppf(0.975)),
    'prior': prior,
    'posterior': posterior
}

return results

# Our trading strategy: 14 wins out of 20 trades
wins = 14
total = 20

# Analyze with uniform prior (no prior belief)
results = bayesian_win_rate_analysis(wins, total)

print("=" * 60)
print("BAYESIAN ANALYSIS: Trading Strategy Win Rate")
print("=" * 60)
print(f"\nObserved: {wins} wins out of {total} trades ({100*wins/total:.0f}% win rate)")
print(f"\nPosterior Analysis:")
print(f"  Expected win rate: {results['posterior_mean']:.1%}")
print(f"  Standard deviation: {results['posterior_std']:.1%}")
print(f"  95% Credible Interval: [{results['hdi_95'][0]:.1%}, {results['hdi_95'][1]:.1%}]")
print(f"\nKey Probabilities:")
print(f"  P(true win rate > 50%) = {results['prob_above_50']:.1%}")
print(f"  P(true win rate > 60%) = {results['prob_above_60']:.1%}")

```

```

In [ ]: # Visualize Prior and Posterior
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Plot range
x = np.linspace(0, 1, 1000)

# Prior vs Posterior
ax = axes[0]
ax.plot(x, results['prior'].pdf(x), 'orange', linewidth=2, label='Prior (Uniform)')
ax.plot(x, results['posterior'].pdf(x), 'blue', linewidth=2, label='Posterior (Bayesian Update)')
ax.axvline(0.5, color='red', linestyle='--', label='Break-even (50%)')
ax.fill_between(x[x > 0.5], results['posterior'].pdf(x[x > 0.5]),
                alpha=0.3, color='green', label=f'P(>50%) = {results["prob_above_50"]:.1%}')
ax.set_xlabel('Win Rate', fontsize=12)
ax.set_ylabel('Probability Density', fontsize=12)
ax.set_title('Bayesian Update: Prior → Posterior', fontsize=14, fontweight='bold')
ax.legend()
ax.set_xlim(0, 1)

# How data changes belief
ax = axes[1]
# Show posteriors with different amounts of data
for n_trades, n_wins in [(5, 4), (20, 14), (100, 70), (500, 350)]:
    posterior = stats.beta(1 + n_wins, 1 + n_trades - n_wins)
    ax.plot(x, posterior.pdf(x), linewidth=2,
            label=f'{n_trades} trades ({100*n_wins/n_trades:.0f}% win)')

```

```

ax.axvline(0.7, color='red', linestyle='--', alpha=0.5, label='True rate')
ax.set_xlabel('Win Rate', fontsize=12)
ax.set_ylabel('Probability Density', fontsize=12)
ax.set_title('More Data → More Certainty', fontsize=14, fontweight='bold')
ax.legend()
ax.set_xlim(0.3, 1)

plt.tight_layout()
plt.show()

```

Key Insight: Uncertainty Decreases with More Data

Notice how the posterior distribution becomes **narrower** (more certain) as we collect more data, peak stays around the true 70% win rate. This is Bayesian learning in action:

- 5 trades: Very uncertain - win rate could be anywhere from 30% to 95%
- 20 trades: Moderate certainty - likely between 50% and 85%
- 100 trades: Pretty confident - likely between 60% and 78%
- 500 trades: Highly confident - tightly concentrated around 70%

Trading implication: Don't bet big on a strategy with only 20 trades of history. The uncertainty is still high!

4 . The Power of Prior Information

One of Bayesian inference's greatest strengths is incorporating prior knowledge. Let's see how different priors affect our conclusions.

Scenario: Experienced vs Naive Trader

Naive trader: "I have no prior beliefs about my strategy" (Uniform prior)

Experienced trader: "Most trading strategies underperform. I expect a 45% win rate on average and I'm fairly confident about this" (Informative prior centered at 0.45)

```

In [ ]: # Compare different priors
wins, total = 14, 20 # Same observed data

priors = {
    'Uniform (Naive)': (1, 1),           # No prior belief
    'Skeptical': (9, 11),                 # Prior mean = 45%, skeptical of
    'Optimistic': (14, 6),                # Prior mean = 70%, believes in
    'Strong Skeptic': (45, 55)           # Prior mean = 45%, very confident
}

fig, axes = plt.subplots(2, 2, figsize=(14, 10))
x = np.linspace(0, 1, 1000)

for ax, (name, (alpha, beta)) in zip(axes.flatten(), priors.items()):
    # Prior
    prior = stats.beta(alpha, beta)

    # Posterior
    post_alpha = alpha + wins
    post_beta = beta + (total - wins)

```

```

posterior = stats.beta(post_alpha, post_beta)

# Plot
ax.plot(x, prior.pdf(x), 'orange', linewidth=2, label='Prior')
ax.plot(x, posterior.pdf(x), 'blue', linewidth=2, label='Posterior')
ax.axvline(0.5, color='red', linestyle='--', alpha=0.5)
ax.axvline(wins/total, color='green', linestyle=':', label='Observed

# Annotations
prob_above_50 = 1 - posterior.cdf(0.5)
ax.set_title(f'{name}\n{prob_above_50:.1%}', fontweight='bold')
ax.set_xlabel('Win Rate')
ax.set_ylabel('Density')
ax.legend(loc='upper left')
ax.set_xlim(0, 1)

plt.tight_layout()
plt.show()

# Summary table
print("\n" + "="*70)
print("SUMMARY: How Priors Affect Conclusions")
print("="*70)
print(f"{'Prior':<20} {'Prior Mean':>12} {'Post. Mean':>12} {'P(>50%)':>12}")
print("-"*70)
for name, (alpha, beta) in priors.items():
    prior_mean = alpha / (alpha + beta)
    post_alpha = alpha + wins
    post_beta = beta + (total - wins)
    post_mean = post_alpha / (post_alpha + post_beta)
    posterior = stats.beta(post_alpha, post_beta)
    prob_above_50 = 1 - posterior.cdf(0.5)
    print(f"{name:<20} {prior_mean:>12.1%} {post_mean:>12.1%} {prob_above_50:>12.1%}")

```

Key Insights on Priors

- 1 . **Priors matter** when data is limited (2 0 trades)
- 2 . **Skeptical priors** protect against overfitting to noise
- 3 . **Strong priors** require more data to overcome
- 4 . **With enough data**, all reasonable priors converge to similar posteriors

Trading Wisdom: In the absence of strong evidence, skepticism is rational. Most trading strategies Bayesian approach naturally encodes this skepticism.

5 . Bayesian vs Frequentist: Credible Intervals vs Confidence Intervals

This is one of the most important distinctions in statistics.

Frequentist 95 % Confidence Interval

"If we repeated this experiment many times and calculated a 95 % CI each time, 95 % intervals would contain the true parameter."

What it does NOT mean: "There's a 95 % probability the true parameter is in this interval."

Bayesian 95% Credible Interval

"Given our prior and the observed data, there's a 95% probability the true parameter is in the interval."

This IS what traders usually want to know!

```
In [ ]: # Compare Frequentist CI vs Bayesian Credible Interval
wins, total = 14, 20
observed_rate = wins / total

# Frequentist: Normal approximation CI
se = np.sqrt(observed_rate * (1 - observed_rate) / total)
freq_ci = (observed_rate - 1.96 * se, observed_rate + 1.96 * se)

# Bayesian: Credible interval from posterior
posterior = stats.beta(1 + wins, 1 + total - wins)
bayes_ci = (posterior.ppf(0.025), posterior.ppf(0.975))

print("=*60)
print("CREDIBLE vs CONFIDENCE INTERVALS")
print("=*60)
print(f"\nObserved: {wins} wins out of {total} trades")
print(f"\nFrequentist 95% CI: [{freq_ci[0]:.1%}, {freq_ci[1]:.1%}]")
print(f" Interpretation: If we repeated this 100 times, ~95 intervals")
print(f" would contain the true win rate.")
print(f"\nBayesian 95% Credible Interval: [{bayes_ci[0]:.1%}, {bayes_ci[1]:.1%}]")
print(f" Interpretation: There's a 95% probability the true win rate")
print(f" is between {bayes_ci[0]:.1%} and {bayes_ci[1]:.1%}")

# Visualize
fig, ax = plt.subplots(figsize=(12, 5))
x = np.linspace(0.3, 1, 1000)
ax.plot(x, posterior.pdf(x), 'blue', linewidth=2, label='Posterior Distribution')
ax.fill_between(x[(x >= bayes_ci[0]) & (x <= bayes_ci[1])],
                posterior.pdf(x[(x >= bayes_ci[0]) & (x <= bayes_ci[1])]),
                alpha=0.3, color='blue', label='95% Credible Interval')
ax.axvline(observed_rate, color='red', linestyle='--', linewidth=2, label='Observed Win Rate')

# Mark frequentist CI
ax.axvline(freq_ci[0], color='orange', linestyle=':', linewidth=2, label='Frequentist CI Lower Bound')
ax.axvline(freq_ci[1], color='orange', linestyle=':', linewidth=2, label='Frequentist CI Upper Bound')

ax.set_xlabel('Win Rate', fontsize=12)
ax.set_ylabel('Probability Density', fontsize=12)
ax.set_title('Bayesian Credible Interval vs Frequentist Confidence Interval')
ax.legend(loc='upper left')
plt.tight_layout()
plt.show()
```

6 . Monte Carlo Simulation: Trading Under Uncertainty

One of the most powerful applications of Bayesian thinking is simulating future outcomes that account for parameter uncertainty.

Traditional Approach

"Our strategy has a 70% win rate. With 100 trades, we expect 70 wins."

Bayesian Approach

"Our strategy's win rate is uncertain (posterior distribution). Let's simulate 100 trades thousands of times, each time drawing a different win rate from our posterior."

This gives us a **distribution of outcomes** that properly accounts for our uncertainty about the trade rate.

```
In [ ]: def simulate_trading_outcomes(posterior, n_future_trades, n_simulations,
                                         avg_win=100, avg_loss=-80):
    """
    Simulate future trading outcomes accounting for parameter uncertainty

    Parameters:
    -----
    posterior : scipy.stats distribution
        Posterior distribution of win rate
    n_future_trades : int
        Number of future trades to simulate
    n_simulations : int
        Number of Monte Carlo simulations
    avg_win : float
        Average profit on winning trade
    avg_loss : float
        Average loss on losing trade

    Returns:
    -----
    final_pnl : array of final P&L for each simulation
    """
    final_pnl = np.zeros(n_simulations)

    for i in range(n_simulations):
        # Step 1: Draw a win rate from posterior
        true_win_rate = posterior.rvs()

        # Step 2: Simulate trades with this win rate
        wins = np.random.binomial(n_future_trades, true_win_rate)
        losses = n_future_trades - wins

        # Step 3: Calculate P&L
        pnl = wins * avg_win + losses * avg_loss
        final_pnl[i] = pnl

    return final_pnl

# Our posterior from earlier
wins, total = 14, 20
posterior = stats.beta(1 + wins, 1 + total - wins)

# Simulate next 100 trades
n_future_trades = 100
```

```

n_simulations = 10000

# Bayesian simulation (accounts for uncertainty)
bayesian_pnl = simulate_trading_outcomes(posterior, n_future_trades, n_simulations)

# Naive simulation (assumes 70% is the true rate)
naive_pnl = np.zeros(n_simulations)
for i in range(n_simulations):
    wins = np.random.binomial(n_future_trades, 0.70)
    losses = n_future_trades - wins
    naive_pnl[i] = wins * 100 + losses * (-80)

# Compare results
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Bayesian simulation
ax = axes[0]
ax.hist(bayesian_pnl, bins=50, density=True, alpha=0.7, color='blue', label='Bayesian')
ax.axvline(0, color='red', linestyle='--', linewidth=2, label='Break-even')
ax.axvline(np.percentile(bayesian_pnl, 5), color='orange', linestyle=':', label='5th Percentile (VaR)')
ax.axvline(np.percentile(bayesian_pnl, 95), color='orange', linestyle=':', label='95th Percentile (VaR)')
ax.set_xlabel('Total P&L ($)', fontsize=12)
ax.set_ylabel('Density', fontsize=12)
ax.set_title('Bayesian: Accounts for Win Rate Uncertainty', fontsize=12, fontweight='bold')
ax.legend()

# Naive simulation
ax = axes[1]
ax.hist(naive_pnl, bins=50, density=True, alpha=0.7, color='green', label='Naive')
ax.axvline(0, color='red', linestyle='--', linewidth=2, label='Break-even')
ax.axvline(np.percentile(naive_pnl, 5), color='orange', linestyle=':', label='5th Percentile (VaR)')
ax.axvline(np.percentile(naive_pnl, 95), color='orange', linestyle=':', label='95th Percentile (VaR)')
ax.set_xlabel('Total P&L ($)', fontsize=12)
ax.set_ylabel('Density', fontsize=12)
ax.set_title('Naive: Assumes 70% Win Rate is Certain', fontsize=12, fontweight='bold')
ax.legend()

plt.tight_layout()
plt.show()

# Print summary
print("\n" + "="*60)
print("MONTE CARLO SIMULATION COMPARISON")
print("="*60)
print(f"\nSimulating {n_future_trades} future trades ({n_simulations:,} simulations)\n")
print(f"\n{'Metric':<25} {'Bayesian':>15} {'Naive':>15}\n")
print("-"*60)
print(f"{'Expected P&L':<25} ${np.mean(bayesian_pnl):>14,.0f} ${np.mean(naive_pnl):>14,.0f}\n")
print(f"{'Std Dev of P&L':<25} ${np.std(bayesian_pnl):>14,.0f} ${np.std(naive_pnl):>14,.0f}\n")
print(f"{'5th Percentile (VaR)':<25} ${np.percentile(bayesian_pnl, 5):>14,.0f} ${np.percentile(naive_pnl, 5):>14,.0f}\n")
print(f"{'P(Loss)':<25} {np.mean(bayesian_pnl < 0):>14.1%} {np.mean(naive_pnl < 0):>14.1%}\n")
print(f"\nBayesian approach shows MORE uncertainty - this is realistic!")

```

Key Insight: Bayesian Simulations Are More Honest About Uncertainty

Notice that the Bayesian simulation has:

- **Higher standard deviation:** Because we're uncertain about the true win rate
- **Fatter tails:** More extreme outcomes are possible
- **Higher probability of loss:** The naive approach underestimates risk

The naive approach falsely assumes we **know** the true win rate is 70% . The Bayesian approach acknowledges our uncertainty and produces more realistic risk estimates.

This is why Bayesian methods matter for risk management.

7 . Real-World Example: Updating Beliefs About Gold

Let's apply Bayesian thinking to a real commodity trading scenario.

Scenario

You're analyzing whether gold prices tend to rise during periods of high uncertainty (measured by VIX). Your prior belief (from reading market research) is that there's a 65% chance of a positive correlation. You want to test this with data.

You collect data for 30 months:

- In 21 months with rising VIX, gold rose in 16 of them (76%)
- This suggests gold does tend to rise during uncertainty

How should you update your beliefs?

```
In [ ]: # Gold-VIX relationship analysis

# Prior: 65% confident gold rises with VIX, moderate certainty
# This translates to Beta(13, 7) which has mean 0.65
prior_alpha = 13
prior_beta = 7

# Observed data
gold_rose_with_vix = 16
total_high_vix_months = 21

# Bayesian update
post_alpha = prior_alpha + gold_rose_with_vix
post_beta = prior_beta + (total_high_vix_months - gold_rose_with_vix)

prior = stats.beta(prior_alpha, prior_beta)
posterior = stats.beta(post_alpha, post_beta)

# Visualize
fig, ax = plt.subplots(figsize=(12, 6))
x = np.linspace(0, 1, 1000)

ax.plot(x, prior.pdf(x), 'orange', linewidth=2, label=f'Prior (mean={prior.mean:.2f})')
ax.plot(x, posterior.pdf(x), 'blue', linewidth=2, label=f'Posterior (mean={posterior.mean:.2f})')
```

```

ax.axvline(0.5, color='red', linestyle='--', alpha=0.5, label='50% (no re
ax.fill_between(x[x > 0.5], posterior.pdf(x[x > 0.5])), alpha=0.3, color='

ax.set_xlabel('P(Gold rises | VIX rises)', fontsize=12)
ax.set_ylabel('Probability Density', fontsize=12)
ax.set_title('Bayesian Analysis: Does Gold Rise When VIX Rises?', fontsz
ax.legend()
ax.set_xlim(0.3, 1)

plt.tight_layout()
plt.show()

# Print analysis
print("\n" + "="*60)
print("GOLD-VIX RELATIONSHIP: Bayesian Analysis")
print("="*60)
print(f"\nPrior belief: Gold rises with VIX {prior.mean():.0%} of the tim
print(f"Observed: {gold_rose_with_vix}/{total_high_vix_months} months gol
print(f"\nPosterior:")
print(f" Updated belief: Gold rises with VIX {posterior.mean():.1%} of t
print(f" 95% Credible Interval: [{posterior.ppf(0.025):.1%}, {posterior.
print(f" P(positive relationship) = {1 - posterior.cdf(0.5):.1%}")
print(f"\nConclusion: Very strong evidence that gold tends to rise during

```

8 . Summary: When to Use Bayesian Methods

Bayesian Methods Excel When:

| Situation | Why Bayesian Helps |
|-----------------------------------|------------------------------------------------------|
| Limited data | Priors regularize estimates and prevent overfitting |
| Expert knowledge available | Priors encode domain expertise |
| Uncertainty quantification needed | Full posterior distribution, not just point estimate |
| Sequential decisions | Natural framework for updating beliefs |
| Risk management | Proper accounting of parameter uncertainty |
| Regime changes possible | Can incorporate beliefs about structural breaks |

Key Concepts Learned

- 1 . **Probability as belief:** Bayesian probability represents our degree of belief, not just long-run frequencies
- 2 . **Prior → Data → Posterior:** We start with beliefs, update with evidence, get refined beliefs
- 3 . **Credible intervals:** " 95 % probability the parameter is here" (what we actually want)
- 4 . **Uncertainty propagation:** Monte Carlo simulation with parameter uncertainty gives realistic estimates
- 5 . **Prior sensitivity:** With limited data, priors matter; with lots of data, they wash out

Knowledge Check Quiz

Test your understanding with these questions. Answers are in the next cell.

Q 1 : In Bayesian inference, what does the "prior" represent?

- A) The data we've collected
- B) Our beliefs before seeing data
- C) The probability of the data
- D) The final answer

Q 2 : A 95% Bayesian credible interval means:

- A) If we repeated the experiment, 95% of intervals would contain the true value
- B) There's a 95% probability the true value is in this interval
- C) The parameter is exactly in this range
- D) We are 95% confident in our methodology

Q 3 : With more data, Bayesian posteriors typically:

- A) Become wider (more uncertain)
- B) Become narrower (more certain)
- C) Stay the same width
- D) Become bimodal

Q 4 : Why might a skeptical prior be appropriate for evaluating trading strategies?

- A) It makes the math easier
- B) Most trading strategies fail, so skepticism is rational
- C) Skeptical priors always give lower risk estimates
- D) The SEC requires it

Q 5 : In the trading simulation, the Bayesian approach showed higher risk because:

- A) The calculations were wrong
- B) It accounts for uncertainty in the true win rate
- C) Bayesian methods always show more risk
- D) The prior was too pessimistic

```
In [ ]: # Quiz Answers
print("*60)
print("QUIZ ANSWERS")
print("*60)
print("""
Q1: B) Our beliefs before seeing data
The prior P(θ) represents our beliefs about the parameter before
observing any data. It encodes domain knowledge.

Q2: B) There's a 95% probability the true value is in this interval
This is the key difference from frequentist confidence intervals!
Bayesian credible intervals have a direct probability interpretation.
```

- Q3: B) Become narrower (more certain)
As we collect more data, our uncertainty decreases and the posterior concentrates around the true value.
- Q4: B) Most trading strategies fail, so skepticism is rational
Given that most strategies underperform after costs, a skeptical prior encodes this base rate and protects against overfitting to noise.
- Q5: B) It accounts for uncertainty in the true win rate
The naive approach assumes we KNOW the win rate is 70%. The Bayesian approach acknowledges we're uncertain about the true rate, leading to wider outcome distributions and more realistic risk estimates.
")

Exercises

Complete these exercises in the `exercises.ipynb` notebook.

Exercise 1 : Bayes' Theorem Calculator (Easy)

Build a function that calculates posterior probabilities given prior, likelihood, and evidence.

Exercise 2 : Strategy Evaluation (Medium)

You have a strategy with 4 5 wins out of 6 0 trades. Calculate:

- The posterior distribution of the win rate
- The probability the true win rate exceeds 6 0 %
- How many more trades you'd need to be 9 5 % confident the rate exceeds 5 0 %

Exercise 3 : Monte Carlo Risk Analysis (Hard)

Extend the Monte Carlo simulation to include:

- Variable win/loss sizes (also uncertain)
- Transaction costs
- Maximum drawdown analysis

Next Module Preview

In Module 2 : Prior Selection and Market Knowledge, we'll learn:

- How to translate domain expertise into prior distributions
- Conjugate priors for computational efficiency
- Prior predictive checks to validate our assumptions
- Encoding commodity seasonality in priors

Module 1 Complete