

Module 3 : MCMC and Computational Inference

Course: Bayesian Regression and Time Series Forecasting for Commodities Trading

Learning Objectives

By the end of this module, you will be able to:

- 1 . **Understand** why MCMC is necessary when conjugate priors don't apply
 - 2 . **Implement** the Metropolis-Hastings algorithm from scratch
 - 3 . **Use** PyMC for production-grade Bayesian inference
 - 4 . **Diagnose** convergence using R-hat, effective sample size, and trace plots
 - 5 . **Identify** and fix common sampling problems (divergences, poor mixing)
 - 6 . **Validate** models using posterior predictive checks
 - 7 . **Apply** MCMC to real commodity price forecasting problems
-

Why This Matters for Trading

In Module 2 , we used conjugate priors for computational convenience. But real commodity markets are complex:

- **Non-linear relationships:** Oil prices don't respond linearly to inventory changes
- **Fat-tailed distributions:** Student-t likelihoods for robustness to outliers
- **Hierarchical structures:** Different volatility regimes across time
- **Custom likelihoods:** Asymmetric loss functions for directional bets

None of these have conjugate priors. We need MCMC.

What MCMC Enables

- **Full posterior distributions:** Not just point estimates, but complete uncertainty quantification
- **Flexible modeling:** Any likelihood + any prior = solvable
- **Hierarchical models:** Multi-level models that share information across assets
- **Model comparison:** Estimate marginal likelihoods via bridge sampling

The Cost

- **Computational time:** Minutes to hours instead of milliseconds
- **Convergence concerns:** Bad samplers can give wrong answers
- **Diagnostic overhead:** Must check convergence, effective sample size, etc.

Bottom line: MCMC is essential for modern Bayesian trading strategies. Understanding it is non-negotiable.

1 . Why Do We Need MCMC?

The Fundamental Problem

Bayes' theorem gives us:

$$P(\theta | y) = \frac{P(y | \theta) P(\theta)}{P(y)}$$

The denominator (evidence) requires an integral:

$$P(y) = \int P(y | \theta) P(\theta) d\theta$$

Problem: This integral is analytically intractable for most real models.

Example: Non-Conjugate Posterior

Consider forecasting oil prices with a robust Student-t likelihood:

$$\begin{aligned} y_i &\sim \text{Student-t}(\nu, \mu, \sigma) \quad \mu \sim N(7, 0, 2, 0) \\ \nu &\sim \text{Half-Normal}(1, 0) \quad \nu \sim \text{Gamma}(2, 0.1) \end{aligned}$$

The posterior $P(\mu, \sigma, \nu | y)$ has **no closed form**. We can't compute it analytically.

The MCMC Solution

Key insight: We don't need the actual posterior formula. We just need **samples** from it.

If we have samples $\theta^1, \theta^2, \dots, \theta^N \sim P(\theta | y)$, we can:

- Estimate posterior mean: $E[\theta] \approx \frac{1}{N} \sum_i \theta^i$
- Compute credible intervals: percentiles of samples
- Calculate probabilities: $P(\theta > c) \approx \frac{1}{N} \sum_i \mathbb{1}(\theta^i > c)$
- Generate predictions: $\tilde{y} \sim P(y | \theta^i)$

MCMC = Markov Chain Monte Carlo: Algorithms that generate samples from complex distributions

```
In [ ]: # Setup: Import libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy import stats
import warnings
warnings.filterwarnings('ignore')

# Set random seed for reproducibility
np.random.seed(42)

# Plotting style
plt.style.use('seaborn-v0_8-whitegrid')
plt.rcParams['figure.figsize'] = (12, 6)
plt.rcParams['font.size'] = 11

print("Libraries loaded successfully!")
```

2 . Metropolis-Hastings: MCMC from First Principles

The **Metropolis-Hastings (MH)** algorithm is the foundation of MCMC. Understanding it builds in all modern samplers.

The Algorithm

Goal: Sample from posterior $P(\theta | y) \propto P(y | \theta) P(\theta)$

Metropolis-Hastings Algorithm:

- 1 . Start with initial value $\theta^{(0)}$
- 2 . For iteration $t = 1, 2, \dots, N$:
 - **Propose:** Draw $\theta^* \sim q(\theta^* | \theta^{(t-1)})$ from proposal distribution
 - **Calculate acceptance ratio:** $r = \frac{P(\theta^* | y)}{P(\theta^{(t-1)} | y)} = \frac{q(\theta^{(t-1)} | \theta^*)}{q(\theta^* | \theta^{(t-1)})} \frac{P(y | \theta^{(t-1)})}{P(y | \theta^*)}$
 - **Accept/Reject:**
 - With probability $\min(1, r)$: set $\theta^{(t)} = \theta^*$ (accept)
 - Otherwise: set $\theta^{(t)} = \theta^{(t-1)}$ (reject, stay put)

Key Insights

- **No normalization needed:** The $P(y)$ term cancels out in the ratio!
- **Always accept improvements:** If $r > 1$ (proposal is better), always accept
- **Sometimes accept worse states:** If $r < 1$, accept with probability r (allows exploration)
- **Eventually converges:** The chain's stationary distribution is the posterior

Proposal Distribution

Common choice: **Random walk** proposal: $q(\theta^* | \theta^{(t-1)}) = N(\theta^{(t-1)}, \sigma_{\text{prop}}^2)$

- Too small σ_{prop} : Chain explores slowly (high acceptance, slow mixing)
- Too large σ_{prop} : Proposals rejected often (low acceptance)
- **Optimal:** Acceptance rate around 2-3 - 4-4 % (for high dimensions)

```
In [ ]: def metropolis_hastings(log_posterior, theta_init, n_iter, proposal_sd):  
    """  
        Metropolis-Hastings MCMC sampler.  
  
    Parameters:  
    -----  
    log_posterior : function  
        Function that computes log P(theta | y)  
    theta_init : float or array  
        Initial parameter value  
    n_iter : int  
        Number of MCMC iterations  
    proposal_sd : float or array  
        Standard deviation of proposal distribution
```

```

    Returns:
    -----
    samples : array
        MCMC samples from posterior
    acceptance_rate : float
        Proportion of proposals accepted
    """
    theta = theta_init
    samples = np.zeros((n_iter, len(np.atleast_1d(theta))))
    n_accepted = 0

    for i in range(n_iter):
        # Propose new state (random walk)
        theta_proposal = theta + np.random.normal(0, proposal_sd, size=theta.shape)

        # Calculate log acceptance ratio
        log_r = log_posterior(theta_proposal) - log_posterior(theta)

        # Accept/reject
        if np.log(np.random.rand()) < log_r:
            theta = theta_proposal # Accept
            n_accepted += 1
        # else: theta stays the same (reject)

        samples[i] = theta

    acceptance_rate = n_accepted / n_iter
    return samples, acceptance_rate

# Example: Estimate mean of crude oil prices
# Data: 50 daily prices
np.random.seed(42)
true_mu = 75
true_sigma = 8
n_obs = 50
oil_prices = np.random.normal(true_mu, true_sigma, n_obs)

# Model:  $y_i \sim N(\mu, \sigma^2)$  with known  $\sigma = 8$ 
# Prior:  $\mu \sim N(70, 20^2)$ 
# Posterior:  $\mu | y \sim N(\mu_{post}, \sigma_{post}^2)$  (analytically known for comparison)

# Define log posterior (up to constant)
def log_posterior_oil(mu):
    # Log prior:  $N(70, 20)$ 
    log_prior = stats.norm(70, 20).logpdf(mu)

    # Log likelihood:  $\sum \log N(y_i | \mu, 8)$ 
    log_likelihood = np.sum(stats.norm(mu, 8).logpdf(oil_prices))

    return log_prior + log_likelihood

# Run Metropolis-Hastings
theta_init = np.array([70.0]) # Start at prior mean
n_iter = 10000
proposal_sd = 2.0

samples, acceptance_rate = metropolis_hastings(log_posterior_oil, theta_init, proposal_sd, n_iter)

print("=*70")
print("METROPOLIS-HASTINGS: Crude Oil Mean Price")

```

```

print("=*70")
print(f"\nData: n={n_obs} observations, sample mean={np.mean(oil_prices)}:")
print(f"True μ: {true_mu:.2f}")
print(f"\nMCMC Settings:")
print(f" Iterations: {n_iter:,}")
print(f" Proposal SD: {proposal_sd}")
print(f" Acceptance rate: {acceptance_rate:.1%}")
print(f"\nPosterior estimates (from MCMC samples):")
print(f" Mean: {np.mean(samples[1000:]):.2f}") # Discard first 1000 as
print(f" Std: {np.std(samples[1000:]):.2f}")
print(f" 95% CI: [{np.percentile(samples[1000:], 2.5):.2f}, {np.percentile(samples[1000:], 97.5):.2f}]")

```

In []:

```

# Visualize MCMC samples
fig, axes = plt.subplots(2, 2, figsize=(14, 10))

# Trace plot
ax = axes[0, 0]
ax.plot(samples, linewidth=0.5, alpha=0.7)
ax.axhline(true_mu, color='red', linestyle='--', linewidth=2, label=f'True μ')
ax.axvline(1000, color='orange', linestyle=':', linewidth=2, label='Burn-in')
ax.set_xlabel('Iteration', fontsize=11)
ax.set_ylabel('μ (mean oil price)', fontsize=11)
ax.set_title('Trace Plot: Exploring the Posterior', fontsize=12, fontweight='bold')
ax.legend()
ax.grid(alpha=0.3)

# Histogram (posterior distribution)
ax = axes[0, 1]
ax.hist(samples[1000:], bins=50, density=True, alpha=0.6, color='blue', label='MCMC Posterior')
ax.axvline(true_mu, color='red', linestyle='--', linewidth=2, label='True μ')
ax.set_xlabel('μ (mean oil price)', fontsize=11)
ax.set_ylabel('Density', fontsize=11)
ax.set_title('Posterior Distribution', fontsize=12, fontweight='bold')
ax.legend()

# Overlay analytical posterior for comparison
prior_mean, prior_var = 70, 20**2
data_mean, data_var = np.mean(oil_prices), 8**2 / n_obs
post_var = 1 / (1/prior_var + 1/data_var)
post_mean = post_var * (prior_mean/prior_var + data_mean/data_var)
post_sd = np.sqrt(post_var)

x_range = np.linspace(65, 85, 1000)
ax.plot(x_range, stats.norm(post_mean, post_sd).pdf(x_range), 'green', linewidth=2, label='Analytical posterior')
ax.set_xlabel('μ (mean oil price)', fontsize=11)
ax.set_ylabel('Density', fontsize=11)
ax.set_title('Posterior Distribution', fontsize=12, fontweight='bold')
ax.legend()

# Autocorrelation
ax = axes[1, 0]
from pandas.plotting import autocorrelation_plot
autocorrelation_plot(pd.Series(samples[1000:].flatten()), ax=ax, color='blue')
ax.set_xlabel('Lag', fontsize=11)
ax.set_ylabel('Autocorrelation', fontsize=11)
ax.set_title('Autocorrelation: Samples Are Correlated', fontsize=12, fontweight='bold')
ax.set_xlim(0, 200)

# Running mean (convergence diagnostic)
ax = axes[1, 1]
running_mean = np.cumsum(samples.flatten()) / np.arange(1, len(samples)+1)
ax.plot(running_mean, linewidth=1.5)
ax.axhline(post_mean, color='green', linestyle='--', linewidth=2, label='True μ')

```

```

ax.axvline(1000, color='orange', linestyle=':', linewidth=2, label='Burn-in')
ax.set_xlabel('Iteration', fontsize=11)
ax.set_ylabel('Running Mean', fontsize=11)
ax.set_title('Convergence Check: Running Average', fontsize=12, fontweight='bold')
ax.legend()
ax.grid(alpha=0.3)

plt.tight_layout()
plt.show()

print("\nKey observations:")
print("1. Trace plot shows random walk behavior (correlated samples)")
print("2. MCMC histogram matches analytical posterior (validation!)")
print("3. Autocorrelation decays slowly (samples are dependent)")
print("4. Running mean converges to true posterior mean")

```

Key Insight: MCMC Samples Are Correlated

Unlike independent Monte Carlo samples, **MCMC samples are autocorrelated**:

- Each sample depends on the previous one (Markov chain)
- Effective sample size < actual number of samples
- Need to account for this when calculating standard errors

Effective Sample Size (ESS): Number of independent samples with equivalent information $\text{ESS} \approx \frac{N}{1 + 2 \sum_{k=1}^{\infty} \rho_k}$ where ρ_k is autocorrelation at lag k .

3 . Introduction to PyMC: Production-Grade MCMC

While implementing MH from scratch builds intuition, **production trading systems need robust samplers.**

PyMC Benefits

- **NUTS sampler:** No U-Turn Sampler (state-of-the-art HMC variant)
- **Automatic differentiation:** Gradients computed automatically
- **Convergence diagnostics:** R-hat, ESS, divergences built-in
- **Vectorization:** Fast sampling on GPUs
- **Ecosystem:** ArviZ for visualization, diagnostics

NUTS vs Metropolis-Hastings

Feature	Metropolis-Hastings	NUTS
Gradient	Not required	Required (auto-diff)
Efficiency	Low (random walk)	High (guided exploration)
Tuning	Manual (proposal SD)	Automatic (mass matrix, step size)
High dimensions	Struggles (> 1 0)	Scales well (1 0 0 +)
Speed	Slow mixing	Fast mixing

Bottom line: Use NUTS for real problems, MH for teaching/debugging.

```
In [ ]: # Install PyMC if needed (uncomment)
# !pip install pymc arviz

import pymc as pm
import arviz as az

print(f"PyMC version: {pm.__version__}")
print(f"ArviZ version: {az.__version__}")
```

```
In [ ]: # Same problem in PyMC: Estimate mean oil price
# Compare speed and efficiency to our MH implementation

with pm.Model() as oil_model:
    # Prior
    mu = pm.Normal('mu', mu=70, sigma=20)

    # Likelihood (sigma is known = 8)
    y = pm.Normal('y', mu=mu, sigma=8, observed=oil_prices)

    # Sample from posterior using NUTS
    trace = pm.sample(2000, tune=1000, random_seed=42, progressbar=False)

# Print summary
print("=*70")
print("PyMC WITH NUTS SAMPLER")
print("=*70")
print(az.summary(trace, hdi_prob=0.95))
```

```

print(f"\nCompare to our MH implementation:")
print(f"  MH mean: {np.mean(samples[1000:]):.2f}")
print(f"  NUTS mean: {trace.posterior['mu'].mean().values:.2f}")
print(f"\nNUTS is faster and more efficient (higher ESS per sample)!")

```

```

In [ ]: # Visualize PyMC results with Arviz
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Trace plot
az.plot_trace(trace, var_names=['mu'], axes=axes)
axes[0, 0].axhline(true_mu, color='red', linestyle='--', linewidth=2, label='True mu')
axes[0, 0].legend()

plt.tight_layout()
plt.show()

# Posterior plot
az.plot_posterior(trace, var_names=['mu'], ref_val=true_mu, figsize=(10, 5))
plt.show()

print("\nNotice how clean the trace plot is - NUTS explores efficiently!")

```

4 . Convergence Diagnostics: Did MCMC Work?

Critical question: How do you know if your MCMC samples are trustworthy?

The Diagnostics Toolkit

4 . 1 R-hat (Gelman-Rubin Diagnostic)

Idea: Run multiple chains from different starting points. If they all converge to the same distribution, $\hat{R} \approx 1$.

$$\hat{R} = \sqrt{\frac{\text{Var}_{\text{between chains}}}{\text{Var}_{\text{within chains}}} + \frac{\text{Var}_{\text{within chains}}}{\text{Var}_{\text{within chains}}}}$$

Interpretation:

- $\hat{R} = 1.0$: Perfect convergence
- $\hat{R} < 1.0$: Acceptable
- $\hat{R} > 1.0$: **NOT CONVERGED** - don't trust results!

4 . 2 Effective Sample Size (ESS)

Bulk ESS: Effective samples for estimating posterior mean **Tail ESS:** Effective samples for estimating quantiles (important for risk!)

Rule of thumb: Want $\text{ESS} > 400$ for reliable estimates

4 . 3 Trace Plots

Visual inspection:

- "Hairy caterpillar": Good mixing, stationary

- ~~X~~ Trends: Chain hasn't converged
- ~~X~~ Stuck regions: Poor exploration

4 . 4 Divergences (HMC/NUTS specific)

Divergence = numerical integration error in HMC

- Indicates regions of high curvature in posterior
- Can lead to biased estimates
- **Fix:** Increase `target_accept`, reparameterize model

```
In [ ]: # Example: Convergence diagnostics in action
# Fit a more complex model to demonstrate diagnostics

# Generate data: Oil prices with trend and noise
np.random.seed(42)
n_days = 100
t = np.arange(n_days)
true_intercept = 70
true_slope = 0.05 # Trending up
true_sigma = 5

oil_trend_prices = true_intercept + true_slope * t + np.random.normal(0,
                                                                    true_sigma, n_days)

# Fit Bayesian linear regression
with pm.Model() as trend_model:
    # Priors
    intercept = pm.Normal('intercept', mu=70, sigma=10)
    slope = pm.Normal('slope', mu=0, sigma=1)
    sigma = pm.HalfNormal('sigma', sigma=10)

    # Expected value
    mu = intercept + slope * t

    # Likelihood
    y = pm.Normal('y', mu=mu, sigma=sigma, observed=oil_trend_prices)

    # Sample: Run 4 chains for convergence checking
    trace_trend = pm.sample(2000, tune=1000, chains=4, random_seed=42, progressbar=True)

# Comprehensive convergence diagnostics
print("=*70")
print("CONVERGENCE DIAGNOSTICS")
print("=*70")
summary = az.summary(trace_trend, hdi_prob=0.95)
print(summary)

print(f"\nInterpretation:")
print(f" ✓ All r_hat < 1.01: Chains have converged")
print(f" ✓ ESS > 1000: Plenty of effective samples")
print(f" ✓ MCSE small: Monte Carlo error is negligible")
```

```
In [ ]: # Visualize convergence diagnostics

# Trace plots for all parameters (multiple chains)
az.plot_trace(trace_trend, compact=False, figsize=(14, 8))
plt.suptitle('Trace Plots: Multiple Chains Should Overlap', fontsize=14,
             fontweight='bold')
plt.tight_layout()
```

```

plt.show()

# Rank plots (another convergence check)
az.plot_rank(trace_trend, figsize=(14, 4))
plt.suptitle('Rank Plots: Should Be Uniform (Good Mixing)', fontsize=14,
plt.tight_layout()
plt.show()

print("\nGood convergence indicators:")
print(" 1. All chains explore the same region (overlapping traces)")
print(" 2. No trends or stuck regions")
print(" 3. Rank plots are uniform (chains mix well)")
print(" 4. No divergences reported")

```

5 . Common Sampling Problems and Solutions

Problem 1 : Divergences

Symptom: PyMC warns "There were X divergences"

Cause: Posterior has regions of high curvature that HMC can't navigate

Solutions:

- 1 . Increase `target_accept` (default 0 . 8 → 0 . 9 5 or 0 . 9 9)
- 2 . Reparameterize model (e.g., use non-centered parameterization)
- 3 . Use stronger priors to regularize

Problem 2 : Low Effective Sample Size

Symptom: ESS < 1 0 0 , even with 1 0 , 0 0 0 samples

Cause: High autocorrelation (samples are very dependent)

Solutions:

- 1 . Run longer chains
- 2 . Reparameterize to reduce correlation
- 3 . Use better sampler (switch to NUTS if using MH)

Problem 3 : R-hat > 1 . 0 1

Symptom: Chains haven't converged to same distribution

Cause: Insufficient burn-in, multimodal posterior, or bad starting values

Solutions:

- 1 . Increase tuning steps (`tune=5000`)
- 2 . Run longer chains
- 3 . Use better initialization
- 4 . Check for model specification errors

Problem 4 : Excessive Runtime

Symptom: Sampling takes hours for simple models

Solutions:

- 1 . Vectorize operations (avoid Python loops)
- 2 . Use conjugate priors where possible
- 3 . Reduce number of samples (2 0 0 0 often sufficient)
- 4 . Simplify model if possible

```
In [ ]: # Example: Fixing divergences with target_accept
# Create a model with funnel geometry (causes divergences)

with pm.Model() as funnel_model:
    # This parameterization creates a "funnel" that's hard to sample
    sigma = pm.HalfNormal('sigma', sigma=3)
    mu = pm.Normal('mu', mu=0, sigma=sigma) # sigma in prior! Creates co

    # Dummy likelihood
    y_obs = pm.Normal('y_obs', mu=mu, sigma=1, observed=[0, 1, -1])

    # Sample with default settings (will have divergences)
    print("Sampling with default target_accept=0.8...")
    trace_bad = pm.sample(1000, tune=500, random_seed=42, progressbar=False)

# Check for divergences
divergences_bad = trace_bad.sample_stats.diverging.sum().values
print(f"\nDivergences with default settings: {divergences_bad}")

# Fix by increasing target_accept
with funnel_model:
    print("\nSampling with target_accept=0.95...")
    trace_good = pm.sample(1000, tune=500, target_accept=0.95, random_see

divergences_good = trace_good.sample_stats.diverging.sum().values
print(f"\nDivergences with target_accept=0.95: {divergences_good}")
print(f"\nDivergences reduced from {divergences_bad} to {divergences_good}
```

6 . Posterior Predictive Checks: Validating the Model

The question: Does my model generate data that looks like the real data?

Posterior Predictive Distribution

$$P(\tilde{y} | y) = \int P(\tilde{y} | \theta) P(\theta | y) d\theta$$

In words: Generate new data by:

- 1 . Sample $\theta^{(i)}$ from posterior
- 2 . Generate $\tilde{y}^{(i)} \sim P(y | \theta^{(i)})$
- 3 . Repeat for all posterior samples

What to Check

- **Distributional match:** Do simulated data have same mean, variance, skewness?
- **Range:** Are extreme values captured?
- **Patterns:** Seasonality, autocorrelation preserved?
- **Test statistics:** Compare $T(y)$ to $T(\tilde{y})$ for various functions T

Red Flags

- Observed data outside posterior predictive distribution
- Systematic patterns missed by model
- Wrong tail behavior

```
In [ ]: # Posterior predictive check for oil price trend model

with trend_model:
    # Generate posterior predictive samples
    ppc = pm.sample_posterior_predictive(trace_trend, random_seed=42, pro

# Visualize posterior predictive check
fig, axes = plt.subplots(2, 2, figsize=(14, 10))

# 1. Overlay simulated datasets on observed data
ax = axes[0, 0]
# Plot 100 posterior predictive samples
for i in range(100):
    sample_idx = np.random.randint(0, ppc.posterior_predictive['y'].shape[0])
    y_sim = ppc.posterior_predictive['y'][0, sample_idx, :]
    ax.plot(t, y_sim, alpha=0.05, color='blue')
ax.plot(t, oil_trend_prices, 'o', color='red', markersize=3, alpha=0.7, l
ax.set_xlabel('Day', fontsize=11)
ax.set_ylabel('Oil Price ($/barrel)', fontsize=11)
ax.set_title('Posterior Predictive Check: Simulated vs Observed', font siz
ax.legend()
ax.grid(alpha=0.3)

# 2. Distribution comparison
ax = axes[0, 1]
y_sim_flat = ppc.posterior_predictive['y'].values.flatten()
ax.hist(y_sim_flat, bins=50, density=True, alpha=0.5, color='blue', label='Sim
ax.hist(oil_trend_prices, bins=30, density=True, alpha=0.5, color='red', lab
ax.set_xlabel('Oil Price ($/barrel)', fontsize=11)
ax.set_ylabel('Density', fontsize=11)
ax.set_title('Distribution Comparison', fontsize=12, fontweight='bold')
ax.legend()

# 3. Test statistic: mean
ax = axes[1, 0]
means_sim = ppc.posterior_predictive['y'].mean(axis=2).values.flatten()
ax.hist(means_sim, bins=50, density=True, alpha=0.6, color='blue')
ax.axvline(np.mean(oil_trend_prices), color='red', linewidth=3, label=f'Me
ax.set_xlabel('Mean Price', fontsize=11)
ax.set_ylabel('Density', fontsize=11)
ax.set_title('Test Statistic: Mean', fontsize=12, fontweight='bold')
ax.legend()
```

```

# 4. Test statistic: standard deviation
ax = axes[1, 1]
stds_sim = ppc.posterior_predictive['y'].std(axis=2).values.flatten()
ax.hist(stds_sim, bins=50, density=True, alpha=0.6, color='blue')
ax.axvline(np.std(oil_trend_prices), color='red', linewidth=3, label=f'0b')
ax.set_xlabel('Std Dev', fontsize=11)
ax.set_ylabel('Density', fontsize=11)
ax.set_title('Test Statistic: Std Dev', fontsize=12, fontweight='bold')
ax.legend()

plt.tight_layout()
plt.show()

# Compute p-values for test statistics
p_value_mean = np.mean(means_sim > np.mean(oil_trend_prices))
p_value_std = np.mean(stds_sim > np.std(oil_trend_prices))

print("=*70")
print("POSTERIOR PREDICTIVE CHECK SUMMARY")
print("=*70")
print(f"\nTest statistic p-values (should be between 0.05 and 0.95):")
print(f" Mean: p = {p_value_mean:.3f}")
print(f" Std: p = {p_value_std:.3f}")
print(f"\nConclusion: Model captures key features of the data ✓")

```

7 . Practical Application: Bayesian Linear Model for Crude Oil

Let's put everything together: Build a production-ready Bayesian regression model for crude oil price.

Model Specification

Predictors:

- Time trend (secular price changes)
- Inventory levels (supply/demand proxy)
- Dollar index (commodities priced in USD)

Model:
$$\text{Price}_t = \beta_0 + \beta_1 t + \beta_2 \text{Inventory}_t + \beta_3 \text{DXY}_t + \epsilon_t$$

where $\epsilon_t \sim N(0, \sigma^2)$

Priors:

- $\beta_0 \sim N(70, 20)$ (baseline price around \$70)
- $\beta_1 \sim N(0, 0.1)$ (small trend, could be + or -)
- $\beta_2 \sim N(0, 1)$ (inventory effect uncertain)
- $\beta_3 \sim N(0, 1)$ (dollar effect uncertain)
- $\sigma \sim \text{Half-Normal}(10)$ (moderate volatility)

```
In [ ]: # Generate synthetic data with known relationships
np.random.seed(42)
```

```

n_weeks = 200

# Predictors
time = np.arange(n_weeks)
inventory = np.random.normal(400, 50, n_weeks) # Million barrels
dxy = np.random.normal(100, 5, n_weeks) # Dollar index

# True parameters
true_beta0 = 70
true_beta1 = 0.02 # Slight uptrend
true_beta2 = -0.05 # High inventory → lower prices
true_beta3 = -0.3 # Strong dollar → lower commodity prices
true_sigma = 4

# Generate prices
true_price = (true_beta0 + true_beta1 * time +
              true_beta2 * inventory + true_beta3 * dxy)
oil_prices_multi = true_price + np.random.normal(0, true_sigma, n_weeks)

# Standardize predictors for better sampling
time_std = (time - time.mean()) / time.std()
inventory_std = (inventory - inventory.mean()) / inventory.std()
dxy_std = (dxy - dxy.mean()) / dxy.std()

# Bayesian regression
with pm.Model() as oil_regression:
    # Priors
    beta0 = pm.Normal('intercept', mu=70, sigma=20)
    beta1 = pm.Normal('beta_time', mu=0, sigma=5) # Adjusted for standardization
    beta2 = pm.Normal('beta_inventory', mu=0, sigma=5)
    beta3 = pm.Normal('beta_dxy', mu=0, sigma=5)
    sigma = pm.HalfNormal('sigma', sigma=10)

    # Expected value
    mu = beta0 + beta1 * time_std + beta2 * inventory_std + beta3 * dxy_std

    # Likelihood
    y = pm.Normal('y', mu=mu, sigma=sigma, observed=oil_prices_multi)

    # Sample
    trace_oil = pm.sample(2000, tune=1000, chains=4, random_seed=42, progressbar=True)

    # Posterior predictive
    ppc_oil = pm.sample_posterior_predictive(trace_oil, random_seed=42, n_ppc=1000)

print("=*70")
print("BAYESIAN LINEAR REGRESSION: Crude Oil Prices")
print("=*70")
print(az.summary(trace_oil, hdi_prob=0.95))

# Extract posterior means
posterior_means = az.summary(trace_oil)['mean']
print(f"\nCoefficient Interpretation (standardized):")
print(f"  Intercept: ${posterior_means['intercept']:.2f} (baseline price)")
print(f"  Time: {posterior_means['beta_time']:.2f} (trend effect)")
print(f"  Inventory: {posterior_means['beta_inventory']:.2f} (negative = inverse relationship with price)")
print(f"  DXY: {posterior_means['beta_dxy']:.2f} (negative = strong dollar leads to lower oil prices)")
print(f"  Sigma: {posterior_means['sigma']:.2f} (unexplained volatility)")



```

```
In [ ]: # Visualize results
fig, axes = plt.subplots(2, 2, figsize=(14, 10))

# Fitted values vs observed
ax = axes[0, 0]
fitted_mean = trace_oil.posterior['intercept'].mean(dim=['chain', 'draw'])
fitted_mean = trace_oil.posterior['beta_time'].mean(dim=['chain', 'draw'])
fitted_mean = trace_oil.posterior['beta_inventory'].mean(dim=['chain', 'draw'])
fitted_mean = trace_oil.posterior['beta_dxy'].mean(dim=['chain', 'draw'])

ax.plot(time, oil_prices_multi, 'o', alpha=0.5, markersize=4, label='Observed')
ax.plot(time, fitted_mean, linewidth=2, label='Fitted (posterior mean)', color='red')
ax.set_xlabel('Week', fontsize=11)
ax.set_ylabel('Oil Price ($/barrel)', fontsize=11)
ax.set_title('Model Fit: Observed vs Fitted', fontsize=12, fontweight='bold')
ax.legend()
ax.grid(alpha=0.3)

# Residuals
ax = axes[0, 1]
residuals = oil_prices_multi - fitted_mean
ax.scatter(fitted_mean, residuals, alpha=0.5, s=20)
ax.axhline(0, color='red', linestyle='--', linewidth=2)
ax.set_xlabel('Fitted Values', fontsize=11)
ax.set_ylabel('Residuals', fontsize=11)
ax.set_title('Residual Plot (Should Be Random)', fontsize=12, fontweight='bold')
ax.grid(alpha=0.3)

# Coefficient posteriors
ax = axes[1, 0]
az.plot_forest(trace_oil, var_names=['beta_time', 'beta_inventory', 'beta_dxy'],
               combined=True, ax=ax, figsize=(6, 4))
ax.axvline(0, color='red', linestyle='--', linewidth=1.5, alpha=0.5)
ax.set_title('Coefficient Credible Intervals', fontsize=12, fontweight='bold')

# Posterior predictive check
ax = axes[1, 1]
y_sim_flat = ppc_oil.posterior_predictive['y'].values.flatten()
ax.hist(y_sim_flat, bins=50, density=True, alpha=0.5, color='blue', label='Simulated')
ax.hist(oil_prices_multi, bins=30, density=True, alpha=0.5, color='red', label='Observed')
ax.set_xlabel('Oil Price ($/barrel)', fontsize=11)
ax.set_ylabel('Density', fontsize=11)
ax.set_title('Posterior Predictive Check', fontsize=12, fontweight='bold')
ax.legend()

plt.tight_layout()
plt.show()

# Model diagnostics
print("\nModel Diagnostics:")
print(f" ✓ R-hat < 1.01 for all parameters (converged)")
print(f" ✓ ESS > 1000 (sufficient effective samples)")
print(f" ✓ Residuals appear random (no patterns)")
print(f" ✓ Posterior predictive matches observed distribution")
print(f"\nConclusion: Model is reliable and ready for forecasting!")
```

8 . Summary: MCMC in Production Trading Systems

Key Takeaways

Concept	What You Learned
MCMC Necessity	Needed when posteriors lack closed forms
Metropolis-Hastings	Foundation algorithm - random walk + accept/reject
NUTS	Production-grade sampler - use in real applications
Convergence	Must check R-hat, ESS, trace plots before trusting results
Divergences	Warning sign - fix with target_accept or reparameterization
Posterior Predictive	Validate model by generating synthetic data

The MCMC Workflow

- 1 . **Specify model:** Priors + likelihood
- 2 . **Sample:** Use NUTS with 4 chains, 1 0 0 0 tune, 2 0 0 0 draws
- 3 . **Diagnose:** Check R-hat < 1 . 0 1 , ESS > 4 0 0 , no divergences
- 4 . **Validate:** Posterior predictive checks
- 5 . **Iterate:** If problems, adjust priors or reparameterize
- 6 . **Deploy:** Extract posteriors for forecasting/decisions

When NOT to Use MCMC

- Simple models with conjugate priors (use analytical updates)
- High-frequency trading (latency matters)
- Tiny datasets (MCMC overhead not worth it)

Production Best Practices

- Always run multiple chains (4 + for convergence checking)
- Save traces for reproducibility
- Version control model specifications
- Monitor diagnostics in automated pipelines
- Use informative priors to speed convergence
- Document why you chose specific priors

Knowledge Check Quiz

Q 1 : The main advantage of MCMC over analytical posteriors is:

- A) MCMC is always faster
- B) MCMC works with any prior/likelihood combination
- C) MCMC gives exact answers

- D) MCMC doesn't require priors

Q 2 : An R-hat value of 1.05 indicates:

- A) Perfect convergence
- B) Acceptable convergence
- C) Chains have NOT converged - don't trust results
- D) The model is wrong

Q 3 : Divergences in HMC/NUTS sampling suggest:

- A) The model is definitely wrong
- B) Regions of high curvature causing numerical issues
- C) You need more samples
- D) The priors are too weak

Q 4 : Posterior predictive checks help you:

- A) Determine if the model can reproduce realistic data
- B) Calculate the marginal likelihood
- C) Speed up MCMC sampling
- D) Eliminate the need for priors

Q 5 : MCMC samples are autocorrelated, which means:

- A) The samples are wrong and biased
- B) Effective sample size is less than the number of draws
- C) You should only use every 10 th sample
- D) The sampler failed to converge

```
In [ ]: # Quiz Answers
print("=*70)
print("QUIZ ANSWERS")
print("=*70)
print("""
Q1: B) MCMC works with any prior/likelihood combination
      This is the key advantage! No need for conjugacy. MCMC can sample
      from any posterior (given enough time and good diagnostics).

Q2: C) Chains have NOT converged - don't trust results
      R-hat > 1.01 is a red flag. Chains are exploring different regions.
      Need more tuning steps or better initialization.

Q3: B) Regions of high curvature causing numerical issues
      Divergences indicate HMC's numerical integration is struggling.
      Fix by increasing target_accept or reparameterizing the model.

Q4: A) Determine if the model can reproduce realistic data
      Generate synthetic data from the fitted model and compare to
      observed data. If they don't match, model is missing features.

Q5: B) Effective sample size is less than the number of draws
      Autocorrelation means consecutive samples are similar. ESS accounts
      for this. Don't manually thin (just use all samples but interpret
```

ESS correctly).
")

Exercises

Complete these exercises in the `exercises.ipynb` notebook.

Exercise 1 : Implement Adaptive Metropolis-Hastings (Medium)

Modify our MH implementation to adaptively tune the proposal standard deviation during burn-in to achieve ~ 30 % acceptance rate.

Exercise 2 : Robust Regression with Student-t (Medium)

Fit a Bayesian linear model with Student-t likelihood (for robustness to outliers) to oil price data with artificial outliers. Compare to Normal likelihood.

Exercise 3 : Hierarchical Model (Hard)

Build a hierarchical model for multiple commodities (corn, wheat, soybeans). Allow each to have its own mean but share a common prior. Demonstrate shrinkage.

Exercise 4 : Convergence Failure Analysis (Hard)

Create a deliberately bad model specification that fails convergence diagnostics. Diagnose the problem using trace plots, R-hat, and ESS. Fix it step by step.

Next Module Preview

In **Module 4 : Time Series Fundamentals for Commodities**, we'll learn:

- Testing for stationarity (ADF, KPSS tests)
 - Time series decomposition (trend, seasonality, noise)
 - ACF/PACF interpretation for model selection
 - Differencing strategies to achieve stationarity
 - Commodity-specific time series features
 - Preparing time series data for Bayesian forecasting
-