

Module 6 : Bayesian Structural Time Series (BSTS)

Learning Objectives

By the end of this module, you will be able to:

- 1 . Understand state-space models and their components
- 2 . Implement local level and local linear trend models
- 3 . Add seasonal components (daily, weekly, annual) to time series
- 4 . Build regression components with external predictors
- 5 . Model time-varying coefficients (dynamic regression)
- 6 . Apply BSTS to gold prices with USD and inflation data
- 7 . Decompose commodity prices into trend, seasonality, and noise

Why This Matters for Trading

Static linear regression assumes relationships don't change. But commodity markets are dynamic:

- **Structural Breaks:** OPEC policy changes, new technology (fracking), regime shifts
- **Time-Varying Relationships:** Oil-USD correlation weakens during crises
- **Complex Seasonality:** Natural gas has annual (winter heating), weekly (storage reports), and daily patterns
- **Trend Changes:** Mean-reversion vs momentum regimes alternate

Bayesian Structural Time Series (BSTS) models decompose prices into interpretable components:

- **Trend:** Long-term direction (is gold in a secular bull market?)
- **Seasonality:** Predictable cycles (harvest pressure on agricultural commodities)
- **Regression:** Impact of fundamentals (how much does USD drive gold today?)
- **Noise:** Unpredictable shocks

This decomposition enables:

- **Better forecasts:** Separate signal from noise
- **Regime detection:** Identify when relationships break down
- **Event impact:** Measure how geopolitical shocks affect trends
- **Seasonality trading:** Exploit recurring patterns with confidence

In short, BSTS gives you X-ray vision into market structure.

```
In [ ]: # Standard imports
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy import stats
import pymc as pm
```

```

import arviz as az
import warnings
warnings.filterwarnings('ignore')

np.random.seed(42)
plt.style.use('seaborn-v0_8-whitegrid')

print(f"PyMC version: {pm.__version__}")
print(f"ArviZ version: {az.__version__}")

```

1 . Introduction to State-Space Models

Theory

A **state-space model** separates observed data from hidden "state" variables:

Observation Equation (what we see):
$$y_t = Z_t \alpha_t + \epsilon_t, \quad \epsilon_t \sim \text{Normal}(0, \sigma_y^2)$$

State Equation (hidden dynamics):
$$\alpha_t = T_t \alpha_{t-1} + R_t \eta_t, \quad \eta_t \sim \text{Normal}(0, \sigma_\alpha^2)$$

Where:

- y_t : Observed price at time t
- α_t : Hidden state (e.g., "true" price level)
- ϵ_t : Observation noise (measurement error, bid-ask bounce)
- η_t : State innovation (real price changes)

Trading Interpretation

- **State α_t** : "Fundamental" or "fair" value
- **Observation noise ϵ_t** : Market microstructure, liquidity shocks
- **State innovation η_t** : New information arrival

Trading strategy: When y_t deviates from α_t (state estimate), mean-revert.

2 . Local Level Model (Random Walk + Noise)

Theory

The simplest state-space model:

$$\begin{aligned} y_t &= \mu_t + \epsilon_t, \quad \epsilon_t \sim \text{Normal}(0, \sigma_y^2) \\ &= \mu_{t-1} + \eta_t, \quad \eta_t \sim \text{Normal}(0, \sigma_\mu^2) \end{aligned}$$

- μ_t : Local level (latent "true price")
- Evolves as a random walk
- $\sigma_\mu^2 / \sigma_y^2$ ratio determines smoothness

Signal-to-Noise Ratio

- High σ_{μ}^2 : Level changes rapidly (trending)
- Low σ_{μ}^2 : Level nearly constant (mean-reverting)
- $\sigma_{\mu}^2 = 0$: Reduces to $y_t = \mu + \epsilon_t$ (white noise around mean)

```
In [ ]: # Generate synthetic data: local level model
np.random.seed(42)
T = 200 # 200 days

# True parameters
true_sigma_mu = 0.5 # State innovation (trend changes)
true_sigma_y = 1.5 # Observation noise
true_mu0 = 50.0 # Initial level

# Generate latent level (random walk)
mu_true = np.zeros(T)
mu_true[0] = true_mu0
for t in range(1, T):
    mu_true[t] = mu_true[t-1] + np.random.normal(0, true_sigma_mu)

# Generate observations
y_obs = mu_true + np.random.normal(0, true_sigma_y, T)

# Visualize
fig, ax = plt.subplots(figsize=(14, 5))
ax.plot(y_obs, 'o-', alpha=0.5, markersize=3, label='Observed Price', color='blue')
ax.plot(mu_true, linewidth=2.5, label='True Latent Level ( $\mu$ )', color='red')
ax.fill_between(range(T), mu_true - 2*true_sigma_y, mu_true + 2*true_sigma_y, alpha=0.2, color='red', label=' $\pm 2\sigma_y$  (Observation Noise)')
ax.set_xlabel('Time (days)', fontsize=12)
ax.set_ylabel('Price', fontsize=12)
ax.set_title('Local Level Model: Latent State vs Observations', fontsize=14)
ax.legend(loc='upper left')
ax.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

print(f"Signal-to-Noise Ratio:  $\sigma_{\mu}^2 / \sigma_y^2 = {true_sigma_mu**2 / true_sigma_y**2}$ ")
print(f"Lower ratio → smoother level (mean-reversion)")
print(f"Higher ratio → rapidly changing level (trending)")
```

```
In [ ]: # Build local level model in PyMC
with pm.Model() as local_level_model:
    # Priors on volatilities
    sigma_mu = pm.HalfNormal('sigma_mu', sigma=2) # State innovation
    sigma_y = pm.HalfNormal('sigma_y', sigma=5) # Observation noise

    # Initial state
    mu_init = pm.Normal('mu_init', mu=50, sigma=10)

    # State innovations
    innovations = pm.Normal('innovations', mu=0, sigma=sigma_mu, shape=T)

    # Build level via cumulative sum
    mu = pm.Deterministic('mu', pm.math.concatenate([[mu_init], mu_init +
```

```

# Observations
y = pm.Normal('y', mu=mu, sigma=sigma_y, observed=y_obs)

# Sample
trace_local_level = pm.sample(2000, tune=1000, return_inferencedata=True,
                               target_accept=0.95, random_seed=42)

print("\nPosterior Summary:")
print(az.summary(trace_local_level, var_names=['sigma_mu', 'sigma_y']))

```

```

In [ ]: # Extract and visualize latent level estimates
mu_posterior = trace_local_level.posterior['mu'].values
mu_mean = mu_posterior.mean(axis=(0, 1))
mu_lower = np.percentile(mu_posterior, 2.5, axis=(0, 1))
mu_upper = np.percentile(mu_posterior, 97.5, axis=(0, 1))

fig, ax = plt.subplots(figsize=(14, 6))
ax.plot(y_obs, 'o', alpha=0.4, markersize=3, label='Observed Price', color='gray')
ax.plot(mu_true, linewidth=2, label='True Level', color='red', linestyle='solid')
ax.plot(mu_mean, linewidth=2.5, label='Estimated Level (Posterior Mean)', color='blue')
ax.fill_between(range(T), mu_lower, mu_upper, alpha=0.3, color='blue',
                label='95% Credible Interval')
ax.set_xlabel('Time (days)', fontsize=12)
ax.set_ylabel('Price', fontsize=12)
ax.set_title('Local Level Model: Filtering Noise from Signal', fontsize=14)
ax.legend(loc='upper left')
ax.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

print("\nTrading Interpretation:")
print("Blue line = 'Fair value' after filtering out noise")
print("When observed price (gray) deviates from blue line → mean-reversion")
print(f"\nEstimated σ_μ: {trace_local_level.posterior['sigma_mu'].values[0]:.2f}")
print(f"Estimated σ_y: {trace_local_level.posterior['sigma_y'].values[0]:.2f}")

```

3 . Local Linear Trend Model

Theory

The local level model has no persistent trend. Add a **slope** component:

$$\begin{aligned}
 y_t &= \mu_t + \epsilon_t, \quad \epsilon_t \sim \text{Normal}(0, \sigma_y^2) \\
 &= \mu_{t-1} + \beta_{t-1} + \eta_t, \quad \eta_t \sim \text{Normal}(0, \sigma_\mu^2) \\
 &\quad \beta_t = \beta_{t-1} + \zeta_t, \quad \zeta_t \sim \text{Normal}(0, \sigma_\beta^2)
 \end{aligned}$$

Where:

- μ_t : Level (intercept)
- β_t : Slope (trend)
- Both evolve as random walks

Trading Interpretation

- $\beta_t > 0$: Uptrend (bullish)

- $\beta_t < 0$: Downtrend (bearish)
- β_t changing sign: Trend reversal
- σ_{β}^2 large: Frequently changing trends (regime switches)

```
In [ ]: # Generate data with time-varying trend
np.random.seed(42)
T_trend = 150

# True parameters
true_sigma_mu_trend = 0.3
true_sigma_beta = 0.05 # Slope changes slowly
true_sigma_y_trend = 1.2

# Initialize
mu_trend_true = np.zeros(T_trend)
beta_trend_true = np.zeros(T_trend)
mu_trend_true[0] = 60.0
beta_trend_true[0] = 0.2 # Initial uptrend

# Evolve states
for t in range(1, T_trend):
    beta_trend_true[t] = beta_trend_true[t-1] + np.random.normal(0, true_
        mu_trend_true[t] = mu_trend_true[t-1] + beta_trend_true[t-1] + np.ran

# Generate observations
y_trend_obs = mu_trend_true + np.random.normal(0, true_sigma_y_trend, T_t

# Visualize
fig, axes = plt.subplots(2, 1, figsize=(14, 8), sharex=True)

# Observed vs Level
axes[0].plot(y_trend_obs, 'o-', alpha=0.5, markersize=3, label='Observed')
axes[0].plot(mu_trend_true, linewidth=2.5, label='True Level ( $\mu$ )', color=
axes[0].set_ylabel('Price', fontsize=12)
axes[0].set_title('Local Linear Trend: Level Component', fontsize=12, fontw
axes[0].legend()
axes[0].grid(True, alpha=0.3)

# Slope over time
axes[1].plot(beta_trend_true, linewidth=2, color='red')
axes[1].axhline(0, color='black', linestyle='--', alpha=0.7, linewidth=1.
    axes[1].fill_between(range(T_trend), 0, beta_trend_true, where=(beta_tren
        alpha=0.3, color='green', label='Uptrend')
    axes[1].fill_between(range(T_trend), 0, beta_trend_true, where=(beta_tren
        alpha=0.3, color='red', label='Downtrend')
    axes[1].set_xlabel('Time (days)', fontsize=12)
    axes[1].set_ylabel('Slope ( $\beta$ )', fontsize=12)
    axes[1].set_title('Time-Varying Trend (Slope Component)', fontsize=12, fo
    axes[1].legend()
    axes[1].grid(True, alpha=0.3)

    plt.tight_layout()
    plt.show()
```

```
In [ ]: # Build local linear trend model
with pm.Model() as llt_model:
    # Priors
    sigma_mu = pm.HalfNormal('sigma_mu', sigma=2)
```

```

sigma_beta = pm.HalfNormal('sigma_beta', sigma=0.5)
sigma_y = pm.HalfNormal('sigma_y', sigma=5)

# Initial states
mu_init = pm.Normal('mu_init', mu=60, sigma=10)
beta_init = pm.Normal('beta_init', mu=0, sigma=1)

# Innovations
innovations_mu = pm.Normal('innovations_mu', mu=0, sigma=sigma_mu, sh
innovations_beta = pm.Normal('innovations_beta', mu=0, sigma=sigma_be

# Build slope (random walk)
beta = pm.Deterministic('beta', pm.math.concatenate(
    [[beta_init], beta_init + pm.math.cumsum(innovations_beta)])
))

# Build level (random walk with drift = slope)
#  $\mu_t = \mu_{t-1} + \beta_{t-1} + \eta_t$ 
# We need to construct this recursively
# Using scan for sequential dependency

def level_step(beta_lag, mu_lag, innovation):
    return mu_lag + beta_lag + innovation

mu_rest, _ = pm.scan(
    fn=level_step,
    sequences=[beta[:-1], innovations_mu],
    outputs_info=[mu_init]
)

mu = pm.Deterministic('mu', pm.math.concatenate([[mu_init], mu_rest]))

# Observations
y = pm.Normal('y', mu=mu, sigma=sigma_y, observed=y_trend_obs)

# Sample
trace_llt = pm.sample(2000, tune=1000, return_inferencedata=True,
                      target_accept=0.95, random_seed=42)

print("\nPosterior Summary:")
print(pm.summary(trace_llt, var_names=['sigma_mu', 'sigma_beta', 'sigma_y']))

```

In []:

```

# Visualize estimates
mu_llt_mean = trace_llt.posterior['mu'].mean(dim=['chain', 'draw']).value
beta_llt_mean = trace_llt.posterior['beta'].mean(dim=['chain', 'draw']).v
beta_llt_lower = np.percentile(trace_llt.posterior['beta'].values, 2.5, a
beta_llt_upper = np.percentile(trace_llt.posterior['beta'].values, 97.5, a

fig, axes = plt.subplots(2, 1, figsize=(14, 8), sharex=True)

# Level
axes[0].plot(y_trend_obs, 'o', alpha=0.3, markersize=3, label='Observed',
axes[0].plot(mu_trend_true, linewidth=2, label='True Level', color='red',
axes[0].plot(mu_llt_mean, linewidth=2.5, label='Estimated Level', color='
axes[0].set_ylabel('Price', fontsize=12)
axes[0].set_title('Estimated Level ( $\mu$ )', fontsize=12, fontweight='bold')
axes[0].legend()
axes[0].grid(True, alpha=0.3)

# Slope

```

```

axes[1].plot(beta_trend_true, linewidth=2, label='True Slope', color='red')
axes[1].plot(beta_llt_mean, linewidth=2.5, label='Estimated Slope', color='blue')
axes[1].fill_between(range(T_trend), beta_llt_lower, beta_llt_upper,
                     alpha=0.3, color='green', label='95% CI')
axes[1].axhline(0, color='black', linestyle='--', alpha=0.7)
axes[1].set_xlabel('Time (days)', fontsize=12)
axes[1].set_ylabel('Slope ( $\beta$ )', fontsize=12)
axes[1].set_title('Estimated Slope (Trend)', fontsize=12, fontweight='bold')
axes[1].legend()
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

# Trading signals
current_slope = beta_llt_mean[-1]
prob_uptrend = np.mean(trace_llt.posterior['beta'].values[:, :, -1] > 0)

print(f"\nCurrent Slope Estimate: {current_slope:.4f}")
print(f"Probability of Uptrend: {prob_uptrend:.2%}")
if prob_uptrend > 0.75:
    print("→ BULLISH: Strong uptrend detected")
elif prob_uptrend < 0.25:
    print("→ BEARISH: Strong downtrend detected")
else:
    print("→ NEUTRAL: Trend direction uncertain")

```

4 . Seasonal Components

Theory

Commodities exhibit strong seasonality:

- **Natural gas:** Winter heating demand
- **Agriculture:** Planting and harvest cycles
- **Power:** Summer cooling load

Add a seasonal component with period \$S\$:

$$\begin{aligned} y_t &= \mu_t + \gamma_t + \epsilon_t \\ \gamma_t &= -\sum_{s=1}^{S-1} \gamma_{s-1} + \omega_t, \quad \omega_t \sim \text{Normal}(0, \sigma_\gamma^2) \end{aligned}$$

The constraint ensures seasonal effects sum to zero over each cycle.

Trading Application

- Identify **high-seasonality months** to increase position size
- **Calendar spreads:** Long winter months, short summer months for nat gas
- **Harvest pressure:** Short agricultural commodities during harvest

```
In [ ]: # Generate data with annual seasonality (12 periods)
np.random.seed(42)
T_seasonal = 120 # 10 years of monthly data
S = 12 # Annual seasonality
```

```

# Trend
trend_seasonal = 50 + 0.05 * np.arange(T_seasonal)

# Seasonal pattern (higher in winter: months 0, 1, 11)
seasonal_pattern = np.array([5, 4, 2, -1, -3, -4, -5, -4, -2, 1, 3, 4])
seasonal_component = np.tile(seasonal_pattern, T_seasonal // S)

# Observations
y_seasonal_obs = trend_seasonal + seasonal_component + np.random.normal(0, 1, len(y))

# Month labels
months = ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec']
month_indices = np.arange(T_seasonal) % 12

# Visualize
fig, axes = plt.subplots(2, 1, figsize=(14, 8), sharex=True)

# Time series
axes[0].plot(y_seasonal_obs, 'o-', alpha=0.6, markersize=3, label='Observation')
axes[0].plot(trend_seasonal, linewidth=2, label='Trend', color='red', linestyle='solid')
axes[0].plot(trend_seasonal + seasonal_component, linewidth=2, label='Trend + Seasonality')
axes[0].set_ylabel('Price', fontsize=12)
axes[0].set_title('Natural Gas Prices with Annual Seasonality', fontsize=14)
axes[0].legend()
axes[0].grid(True, alpha=0.3)

# Seasonal pattern
axes[1].bar(range(12), seasonal_pattern, color='steelblue', alpha=0.7)
axes[1].set_xticks(range(12))
axes[1].set_xticklabels(months)
axes[1].set_ylabel('Seasonal Effect', fontsize=12)
axes[1].set_title('Monthly Seasonal Pattern (Winter Premium)', fontsize=14)
axes[1].axhline(0, color='black', linestyle='--')
axes[1].grid(True, alpha=0.3, axis='y')

plt.tight_layout()
plt.show()

```

In []:

```

# Build model with trend + seasonality
# For simplicity, we'll use a fixed seasonal pattern approach
# (More sophisticated: stochastic seasonality with state evolution)

with pm.Model() as seasonal_model:
    # Trend component (local level)
    sigma_mu = pm.HalfNormal('sigma_mu', sigma=1)
    mu_init = pm.Normal('mu_init', mu=50, sigma=10)
    innovations_mu = pm.Normal('innovations_mu', mu=0, sigma=sigma_mu, shape=S)
    mu = pm.Deterministic('mu', pm.math.concatenate([
        [mu_init], mu_init + pm.math.cumsum(innovations_mu)]))

    # Seasonal component (one parameter per month, sum-to-zero constraint)
    # We'll use a centered parameterization
    seasonal_raw = pm.Normal('seasonal_raw', mu=0, sigma=5, shape=S-1)
    # Last seasonal component ensures sum = 0
    seasonal_effects = pm.Deterministic('seasonal_effects',
                                         pm.math.concatenate([seasonal_raw,
                                         seasonal_raw[-1]]))

    # Map seasonal effects to each time point
    seasonal_component = seasonal_effects[month_indices]

```

```

# Observation noise
sigma_y = pm.HalfNormal('sigma_y', sigma=5)

# Combined mean
mean = mu + seasonal_component

# Likelihood
y = pm.Normal('y', mu=mean, sigma=sigma_y, observed=y_seasonal_obs)

# Sample
trace_seasonal = pm.sample(2000, tune=1000, return_inferencedata=True
                           target_accept=0.95, random_seed=42)

print("\nPosterior Summary:")
print(pm.summary(trace_seasonal, var_names=['sigma_mu', 'sigma_y', 'seaso

```

```

In [ ]: # Extract estimates
mu_seasonal_mean = trace_seasonal.posterior['mu'].mean(dim=['chain', 'draw'])
seasonal_effects_mean = trace_seasonal.posterior['seasonal_effects'].mean(dim=['chain', 'draw'])
seasonal_component_fitted = seasonal_effects_mean[month_indices]

fig, axes = plt.subplots(2, 1, figsize=(14, 8))

# Decomposition
axes[0].plot(y_seasonal_obs, 'o', alpha=0.4, markersize=3, label='Observed')
axes[0].plot(mu_seasonal_mean, linewidth=2, label='Estimated Trend', color='red')
axes[0].plot(mu_seasonal_mean + seasonal_component_fitted, linewidth=2,
            label='Trend + Seasonal', color='blue')
axes[0].set_ylabel('Price', fontsize=12)
axes[0].set_title('Decomposition: Trend + Seasonality', fontsize=12, fontweight='bold')
axes[0].legend()
axes[0].grid(True, alpha=0.3)

# Seasonal effects comparison
x_pos = np.arange(12)
axes[1].bar(x_pos - 0.2, seasonal_pattern, width=0.4, label='True', alpha=0.6)
axes[1].bar(x_pos + 0.2, seasonal_effects_mean, width=0.4, label='Estimated', alpha=0.6)
axes[1].set_xticks(x_pos)
axes[1].set_xticklabels(months)
axes[1].set_ylabel('Seasonal Effect', fontsize=12)
axes[1].set_title('Estimated vs True Seasonal Pattern', fontsize=12, fontweight='bold')
axes[1].axhline(0, color='black', linestyle='--')
axes[1].legend()
axes[1].grid(True, alpha=0.3, axis='y')

plt.tight_layout()
plt.show()

# Trading signals
print("\n==== SEASONAL TRADING STRATEGY ===")
sorted_months = np.argsort(seasonal_effects_mean)[::-1]
print("\nMonths with HIGHEST prices (go long ahead of these):")
for i in range(3):
    month_idx = sorted_months[i]
    print(f" {months[month_idx]}: +${{seasonal_effects_mean[month_idx]}:.2f}")

print("\nMonths with LOWEST prices (go short or avoid):")
for i in range(3):
    month_idx = sorted_months[i]
    print(f" {months[month_idx]}: -${{seasonal_effects_mean[month_idx]}:.2f}")

```

```

month_idx = sorted_months[-(i+1)]
print(f" {months[month_idx]}: ${seasonal_effects_mean[month_idx]:.2f}

```

5 . Dynamic Regression with Time-Varying Coefficients

Theory

In static regression, β is constant. In **dynamic regression**, coefficients evolve:

$$\begin{aligned} y_t &= \alpha_t + \beta_t x_t + \epsilon_t \\ \alpha_t &= \alpha_{t-1} + \eta_\alpha \\ \beta_t &= \beta_{t-1} + \eta_\beta \end{aligned}$$

Trading Application

Gold-USD relationship changes:

- **Normal times:** $\beta < 0$ \$ (inverse relationship)
- **Flight to safety:** $\beta \approx 0$ \$ (decoupling)
- **Currency crisis:** β very negative

Dynamic regression detects these regime changes automatically.

```

In [ ]: # Generate data with time-varying coefficient
np.random.seed(42)
T_dyn = 200

# Predictor: USD index
usd_index = 95 + np.cumsum(np.random.normal(0, 0.3, T_dyn))
usd_standardized = (usd_index - usd_index.mean()) / usd_index.std()

# Time-varying coefficient (becomes more negative over time → stronger inverse)
true_beta_dyn = np.zeros(T_dyn)
true_beta_dyn[0] = -0.5
for t in range(1, T_dyn):
    true_beta_dyn[t] = true_beta_dyn[t-1] + np.random.normal(0, 0.02)

# Gold price
true_alpha_dyn = 1500
gold_price = true_alpha_dyn + true_beta_dyn * usd_standardized * 100 + np.random.normal(0, 10)

# Visualize
fig, axes = plt.subplots(3, 1, figsize=(14, 10), sharex=True)

# Gold price
axes[0].plot(gold_price, linewidth=1.5, color='gold')
axes[0].set_ylabel('Gold Price ($/oz)', fontsize=12)
axes[0].set_title('Gold Prices', fontsize=12, fontweight='bold')
axes[0].grid(True, alpha=0.3)

# USD Index
axes[1].plot(usd_index, linewidth=1.5, color='green')
axes[1].set_ylabel('USD Index', fontsize=12)
axes[1].set_title('USD Index (Predictor)', fontsize=12, fontweight='bold')
axes[1].grid(True, alpha=0.3)

# Time-varying coefficient

```

```

        axes[2].plot(true_beta_dyn, linewidth=2, color='red')
        axes[2].axhline(0, color='black', linestyle='--')
        axes[2].fill_between(range(T_dyn), 0, true_beta_dyn, where=(true_beta_dyn < 0),
                             alpha=0.3, color='red', label='Inverse Relationship')
        axes[2].set_xlabel('Time (days)', fontsize=12)
        axes[2].set_ylabel('β (Gold-USD)', fontsize=12)
        axes[2].set_title('Time-Varying Coefficient: Gold vs USD', fontsize=12, fontweight='bold')
        axes[2].legend()
        axes[2].grid(True, alpha=0.3)

    plt.tight_layout()
    plt.show()

```

In []:

```

# Build dynamic regression model
with pm.Model() as dynamic_reg_model:
    # Time-varying intercept
    sigma_alpha = pm.HalfNormal('sigma_alpha', sigma=10)
    alpha_init = pm.Normal('alpha_init', mu=1500, sigma=100)
    innovations_alpha = pm.Normal('innovations_alpha', mu=0, sigma=sigma_alpha)
    alpha = pm.Deterministic('alpha', pm.math.concatenate(
        [[alpha_init], alpha_init + pm.math.cumsum(innovations_alpha)]))

    # Time-varying coefficient on USD
    sigma_beta = pm.HalfNormal('sigma_beta', sigma=0.5)
    beta_init = pm.Normal('beta_init', mu=-0.5, sigma=1)
    innovations_beta = pm.Normal('innovations_beta', mu=0, sigma=sigma_beta)
    beta = pm.Deterministic('beta', pm.math.concatenate(
        [[beta_init], beta_init + pm.math.cumsum(innovations_beta)]))

    # Observation noise
    sigma_y = pm.HalfNormal('sigma_y', sigma=50)

    # Regression equation
    mu = alpha + beta * usd_standardized * 100

    # Likelihood
    y = pm.Normal('y', mu=mu, sigma=sigma_y, observed=gold_price)

    # Sample
    trace_dyn_reg = pm.sample(2000, tune=1000, return_inferencedata=True,
                              target_accept=0.95, random_seed=42)

print("\nPosterior Summary:")
print(az.summary(trace_dyn_reg, var_names=['sigma_alpha', 'sigma_beta', 'alpha', 'beta']))

```

In []:

```

# Visualize time-varying coefficient
beta_dyn_mean = trace_dyn_reg.posterior['beta'].mean(dim=['chain', 'draw'])
beta_dyn_lower = np.percentile(trace_dyn_reg.posterior['beta'].values, 2.5)
beta_dyn_upper = np.percentile(trace_dyn_reg.posterior['beta'].values, 97.5)

fig, axes = plt.subplots(2, 1, figsize=(14, 8), sharex=True)

# Fitted vs Observed
alpha_dyn_mean = trace_dyn_reg.posterior['alpha'].mean(dim=['chain', 'draw'])
fitted = alpha_dyn_mean + beta_dyn_mean * usd_standardized * 100

axes[0].plot(gold_price, 'o', alpha=0.4, markersize=3, label='Observed',
             color='blue')

```

```

axes[0].plot(fitted, linewidth=2, label='Fitted', color='gold')
axes[0].set_ylabel('Gold Price ($/oz)', fontsize=12)
axes[0].set_title('Dynamic Regression Fit', fontsize=12, fontweight='bold')
axes[0].legend()
axes[0].grid(True, alpha=0.3)

# Time-varying coefficient
axes[1].plot(true_beta_dyn, linewidth=2, label='True β', color='red', lin
axes[1].plot(beta_dyn_mean, linewidth=2.5, label='Estimated β', color='bl
axes[1].fill_between(range(T_dyn), beta_dyn_lower, beta_dyn_upper,
                     alpha=0.3, color='blue', label='95% CI')
axes[1].axhline(0, color='black', linestyle='--', alpha=0.7)
axes[1].set_xlabel('Time (days)', fontsize=12)
axes[1].set_ylabel('β (Gold-USD)', fontsize=12)
axes[1].set_title('Time-Varying Coefficient: Gold vs USD', fontsize=12, f
axes[1].legend()
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

# Trading signals
current_beta = beta_dyn_mean[-1]
recent_beta_change = beta_dyn_mean[-1] - beta_dyn_mean[-20]

print(f"\n==== GOLD-USD RELATIONSHIP ===")
print(f"Current β: {current_beta:.3f}")
print(f"20-day change in β: {recent_beta_change:.3f}")
if current_beta < -0.5:
    print("→ STRONG INVERSE: USD strength heavily pressures gold")
    print("   Strategy: Short gold when USD rallies")
elif current_beta > -0.2:
    print("→ WEAK INVERSE or DECOUPLED: Relationship breakdown")
    print("   Strategy: USD less reliable predictor, watch other drivers")
else:
    print("→ MODERATE INVERSE: Typical negative correlation")
    print("   Strategy: Use USD as hedge factor")

```

6 . Practical Application: Gold Prices with USD and Infl.

Trading Context

Gold is driven by:

- 1 . **USD strength**: Inverse (commodities priced in USD)
- 2 . **Inflation expectations**: Positive (inflation hedge)
- 3 . **Real rates**: Inverse (opportunity cost)

We'll build a full BSTS model with:

- Local linear trend
- Dynamic regression on USD and inflation
- Student-t likelihood (robust to outliers)

```
In [ ]: # Generate realistic gold price data
np.random.seed(42)
T_gold = 250 # ~1 year of daily data
```

```

# Predictors
usd = 95 + np.cumsum(np.random.normal(0, 0.2, T_gold)) # USD index
inflation = 2.5 + np.cumsum(np.random.normal(0, 0.05, T_gold)) # Inflation

# Standardize
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
predictors = scaler.fit_transform(np.column_stack([usd, inflation]))
usd_std = predictors[:, 0]
inflation_std = predictors[:, 1]

# Time-varying coefficients
beta_usd_true = np.zeros(T_gold)
beta_inflation_true = np.zeros(T_gold)
beta_usd_true[0] = -15 # Negative: strong USD → lower gold
beta_inflation_true[0] = 25 # Positive: inflation → higher gold

for t in range(1, T_gold):
    beta_usd_true[t] = beta_usd_true[t-1] + np.random.normal(0, 0.5)
    beta_inflation_true[t] = beta_inflation_true[t-1] + np.random.normal(0, 0.5)

# Trend (local linear trend)
mu_gold = np.zeros(T_gold)
beta_trend = np.zeros(T_gold)
mu_gold[0] = 1800
beta_trend[0] = 0.3

for t in range(1, T_gold):
    beta_trend[t] = beta_trend[t-1] + np.random.normal(0, 0.05)
    mu_gold[t] = mu_gold[t-1] + beta_trend[t-1] + np.random.normal(0, 2)

# Gold price
gold_full = (mu_gold +
              beta_usd_true * usd_std +
              beta_inflation_true * inflation_std +
              np.random.standard_t(df=5, size=T_gold) * 8) # Fat tails

# Visualize
fig, axes = plt.subplots(4, 1, figsize=(14, 12), sharex=True)

axes[0].plot(gold_full, linewidth=1.5, color='gold')
axes[0].set_ylabel('Gold ($/oz)', fontsize=11)
axes[0].set_title('Gold Prices', fontsize=12, fontweight='bold')
axes[0].grid(True, alpha=0.3)

axes[1].plot(usd, linewidth=1.5, color='green')
axes[1].set_ylabel('USD Index', fontsize=11)
axes[1].set_title('USD Index', fontsize=12, fontweight='bold')
axes[1].grid(True, alpha=0.3)

axes[2].plot(inflation, linewidth=1.5, color='red')
axes[2].set_ylabel('Inflation (%)', fontsize=11)
axes[2].set_title('Inflation Expectations', fontsize=12, fontweight='bold')
axes[2].grid(True, alpha=0.3)

axes[3].plot(beta_usd_true, linewidth=2, label='β (USD)', color='green')
axes[3].plot(beta_inflation_true, linewidth=2, label='β (Inflation)', color='red')
axes[3].axhline(0, color='black', linestyle='--')
axes[3].set_xlabel('Time (days)', fontsize=11)

```

```

axes[3].set_ylabel('Coefficient', fontsize=11)
axes[3].set_title('Time-Varying Coefficients (True)', fontsize=12, fontweight='bold')
axes[3].legend()
axes[3].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```

In []:

```

# Build full BSTS model
# This is computationally intensive, so we'll use a simplified version

with pm.Model() as bststs_full:
    # Trend component (local level only for speed)
    sigma_mu = pm.HalfNormal('sigma_mu', sigma=5)
    mu_init = pm.Normal('mu_init', mu=1800, sigma=100)
    innovations_mu = pm.Normal('innovations_mu', mu=0, sigma=sigma_mu, shape=100)
    mu = pm.Deterministic('mu', pm.math.concatenate(
        [[mu_init], mu_init + pm.math.cumsum(innovations_mu)])
    ))

    # Time-varying USD coefficient
    sigma_beta_usd = pm.HalfNormal('sigma_beta_usd', sigma=2)
    beta_usd_init = pm.Normal('beta_usd_init', mu=-15, sigma=10)
    innovations_beta_usd = pm.Normal('innovations_beta_usd', mu=0, sigma=sigma_beta_usd)
    beta_usd = pm.Deterministic('beta_usd', pm.math.concatenate(
        [[beta_usd_init], beta_usd_init + pm.math.cumsum(innovations_beta_usd)])
    ))

    # Time-varying inflation coefficient
    sigma_beta_inf = pm.HalfNormal('sigma_beta_inf', sigma=2)
    beta_inf_init = pm.Normal('beta_inf_init', mu=25, sigma=10)
    innovations_beta_inf = pm.Normal('innovations_beta_inf', mu=0, sigma=sigma_beta_inf)
    beta_inf = pm.Deterministic('beta_inf', pm.math.concatenate(
        [[beta_inf_init], beta_inf_init + pm.math.cumsum(innovations_beta_inf)])
    ))

    # Robust likelihood (Student-t)
    nu = pm.Gamma('nu', alpha=2, beta=0.1)
    sigma_y = pm.HalfNormal('sigma_y', sigma=20)

    # Combined mean
    mean = mu + beta_usd * usd_std + beta_inf * inflation_std

    # Likelihood
    y = pm.StudentT('y', nu=nu, mu=mean, sigma=sigma_y, observed=gold_full)

    # Sample
    trace_bsts_full = pm.sample(1500, tune=1000, return_inferencedata=True,
                                 target_accept=0.95, random_seed=42, chains=4)

print("\nPosterior Summary:")
print(pm.summary(trace_bsts_full, var_names=['sigma_mu', 'sigma_beta_usd',
                                              'sigma_y', 'nu']))

```

In []:

```

# Decompose gold prices
mu_est = trace_bsts_full.posterior['mu'].mean(dim=['chain', 'draw']).values
beta_usd_est = trace_bsts_full.posterior['beta_usd'].mean(dim=['chain', 'draw']).values
beta_inf_est = trace_bsts_full.posterior['beta_inf'].mean(dim=['chain', 'draw']).values

```

```

# Components
usd_effect = beta_usd_est * usd_std
inflation_effect = beta_inf_est * inflation_std
fitted = mu_est + usd_effect + inflation_effect

fig, axes = plt.subplots(3, 1, figsize=(14, 10), sharex=True)

# Observed vs Fitted
axes[0].plot(gold_full, 'o', alpha=0.4, markersize=3, label='Observed', color='blue')
axes[0].plot(fitted, linewidth=2, label='Fitted (BSTS)', color='gold')
axes[0].set_ylabel('Gold Price ($/oz)', fontsize=11)
axes[0].set_title('BSTS Model: Fitted vs Observed', fontsize=12, fontweight='bold')
axes[0].legend()
axes[0].grid(True, alpha=0.3)

# Component decomposition
axes[1].plot(mu_est, linewidth=2, label='Trend ( $\mu$ )', color='black')
axes[1].plot(mu_est + usd_effect, linewidth=2, label='Trend + USD', color='red')
axes[1].plot(fitted, linewidth=2, label='Trend + USD + Inflation', color='blue')
axes[1].set_ylabel('Price Components ($/oz)', fontsize=11)
axes[1].set_title('Price Decomposition', fontsize=12, fontweight='bold')
axes[1].legend()
axes[1].grid(True, alpha=0.3)

# Time-varying coefficients
axes[2].plot(beta_usd_true, linewidth=2, label='True  $\beta$  (USD)', color='grey')
axes[2].plot(beta_usd_est, linewidth=2, label='Est  $\beta$  (USD)', color='green')
axes[2].plot(beta_inflation_true, linewidth=2, label='True  $\beta$  (Inflation)')
axes[2].plot(beta_inf_est, linewidth=2, label='Est  $\beta$  (Inflation)', color='blue')
axes[2].axhline(0, color='black', linestyle='--', alpha=0.5)
axes[2].set_xlabel('Time (days)', fontsize=11)
axes[2].set_ylabel('Coefficient', fontsize=11)
axes[2].set_title('Time-Varying Coefficients', fontsize=12, fontweight='bold')
axes[2].legend(loc='best', fontsize=9)
axes[2].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

print(f"\nModel RMSE: ${np.sqrt(np.mean((gold_full - fitted)**2)):.2f}")

```

In []: # Trading signals based on decomposition

```

current_trend = mu_est[-1]
current_beta_usd = beta_usd_est[-1]
current_beta_inf = beta_inf_est[-1]
current_usd_effect = usd_effect[-1]
current_inf_effect = inflation_effect[-1]
current_fair_value = fitted[-1]
current_price = gold_full[-1]
mispricing = current_price - current_fair_value

print("\n" + "="*60)
print("GOLD TRADING DASHBOARD")
print("="*60)

print(f"\nCurrent Price: ${current_price:.2f}")
print(f"Fair Value (BSTS): ${current_fair_value:.2f}")
print(f"Mispricing: ${mispricing:+.2f} ({mispricing/current_fair_value*100:.2f}% mispricing)")

if abs(mispricing) > 20:
    print("SELL SIGNAL")
else:
    print("BUY SIGNAL")

```

```

if mispricing > 0:
    print("→ OVERVALUED: Consider short position or tighten stops on")
else:
    print("→ UNDERRATED: Consider long position or add to existing long")
else:
    print("→ FAIRLY VALUED: No mispricing signal")

print(f"\n--- PRICE DECOMPOSITION ---")
print(f"Trend component: ${current_trend:.2f}")
print(f"USD effect: ${current_usd_effect:+.2f} (β = {current_beta_usd:.2f})")
print(f"Inflation effect: ${current_inf_effect:+.2f} (β = {current_beta_inf:.2f})")

print(f"\n--- FACTOR SENSITIVITIES ---")
print(f"β (USD): {current_beta_usd:.2f}")
if current_beta_usd < -10:
    print("→ STRONG INVERSE: USD rallies hurt gold significantly")
    print("    Strategy: Hedge gold positions with long USD")
elif current_beta_usd > -5:
    print("→ WEAK INVERSE: Gold-USD correlation breaking down")
    print("    Strategy: USD less reliable hedge, watch other factors")

print(f"\nβ (Inflation): {current_beta_inf:.2f}")
if current_beta_inf > 20:
    print("→ STRONG INFLATION HEDGE: Rising inflation supports gold")
    print("    Strategy: Increase gold allocation in inflationary environment")
elif current_beta_inf < 10:
    print("→ WEAK INFLATION HEDGE: Gold not responding to inflation")
    print("    Strategy: Consider other inflation hedges (TIPS, commodities)")

print("\n" + "="*60)

```

Knowledge Check Quiz

Question 1

What is the key advantage of state-space models over standard regression?

- A) Faster computation
- B) Separate latent "true" state from noisy observations
- C) Require fewer data points
- D) Always more accurate

Answer: B - State-space models explicitly model hidden states (e.g., "fair value") that evolve over time, separating signal from noise.

Question 2

In a local level model, what does the ratio $\sigma^2_\mu / \sigma^2_y$ control?

- A) Model accuracy
- B) Smoothness of the latent level
- C) Number of parameters
- D) Forecast horizon

Answer: B - Higher ratio → level changes rapidly (trending). Lower ratio → level nearly constant (reverting).

Question 3

Why are seasonal components important for commodity trading?

- A) They increase model complexity
- B) They capture predictable price patterns (harvest, weather)
- C) They are required by PyMC
- D) They replace the need for fundamentals

Answer: B - Commodities have strong seasonal patterns (e.g., natural gas winter demand, agricultural harvest pressure) that create tradable opportunities.

Question 4

What does a time-varying coefficient (dynamic regression) detect?

- A) Data errors
- B) Regime changes in relationships between variables
- C) Missing data
- D) Outliers

Answer: B - Dynamic regression allows coefficients to evolve, detecting when relationships strengthen, weaken, or reverse (e.g., Gold-USD correlation during crises).

Question 5

In the gold BSTS model, what does it mean if β (USD) becomes less negative over time?

- A) USD is getting stronger
- B) The inverse Gold-USD relationship is weakening
- C) Gold is becoming more volatile
- D) Model is overfitting

Answer: B - A coefficient moving toward zero indicates the relationship is weakening (decoupling), possibly due to regime change or other dominant factors.

Exercises

Exercise 1 : Crude Oil Seasonality

Build a BSTS model for crude oil with:

- Local linear trend

- Weekly seasonality ($S= 5$ for trading days)
- Compare model with vs without seasonality using WAIC
- Identify which day of the week has highest/lowest prices

Exercise 2 : Natural Gas Multi-Seasonality

Natural gas has multiple seasonal patterns:

- Annual ($S= 12$ months): Winter heating
- Weekly ($S= 7$ days): Storage report Thursdays

Build a model with both. Which is more important?

Exercise 3 : Corn Prices with Weather

Generate synthetic corn data with:

- Annual harvest seasonality (low prices in fall)
- Dynamic regression on rainfall (time-varying coefficient)
- Simulate drought year where β (rainfall) spikes
- Detect the regime change automatically

Exercise 4 : Regime Detection

Using the gold BSTS model:

- 1 . Define a regime as "when β (USD) < -2 0 "
- 2 . Calculate P(regime) at each time point
- 3 . Visualize regime probabilities over time
- 4 . Design a trading rule: increase position size when $P(\text{high sensitivity regime}) > 0.8$

Exercise 5 : Forecast Decomposition

For the gold model:

- 1 . Forecast 30 days ahead
 - 2 . Decompose forecast into: trend + USD effect + inflation effect
 - 3 . Scenario analysis:
 - If USD rises 2 %, how much does gold forecast drop?
 - If inflation rises 0.5 %, how much does gold forecast rise?
 - 4 . Calculate forecast credible intervals
-

Summary

In this module, you learned:

- 1 . **State-Space Models:** Separate latent states from noisy observations
- 2 . **Local Level Model:** Random walk for slowly changing "fair value"

- 3 . **Local Linear Trend:** Add time-varying slope for trend detection
- 4 . **Seasonal Components:** Model predictable cycles (harvest, weather)
- 5 . **Dynamic Regression:** Time-varying coefficients detect regime changes
- 6 . **Gold Application:** Full BSTS with trend, USD, and inflation factors

Key Takeaways for Trading

- **Decomposition:** Break prices into trend + seasonal + fundamental + noise
- **Fair Value:** Use latent level estimate for mean-reversion trades
- **Regime Detection:** Time-varying coefficients automatically identify structural breaks
- **Seasonality:** Exploit recurring patterns with statistical confidence
- **Factor Sensitivity:** Track how responsive commodities are to drivers (USD, inflation)
- **Robust Estimation:** Student-t likelihood handles price shocks without bias

BSTS models provide a principled framework for understanding commodity price dynamics in non-stationary, regime-switching markets.

Preview of Next Module

Module 7 : Hierarchical Models for Multiple Commodities

So far, we've modeled one commodity at a time. But markets are interconnected:

- **Energy complex:** WTI, Brent, Natural Gas share supply shocks
- **Grains:** Corn, Wheat, Soybeans compete for farmland
- **Metals:** Gold, Silver, Copper respond to economic cycles

Hierarchical models pool information across related commodities:

- Estimate **group-level parameters** (e.g., average energy seasonality)
- Allow **commodity-specific deviations** (Brent vs WTI spreads)
- **Shrinkage:** Regularize estimates for thinly-traded commodities

You'll learn:

- Partial pooling vs complete pooling vs no pooling
- Varying intercepts and slopes
- Cross-commodity correlation structures
- Pairs trading with hierarchical models

See you in Module 7 !